



# Parallelizing the Naughty Dog engine using fibers

**Christian Gyrling**

Lead Programmer @ Naughty Dog

GAME DEVELOPERS CONFERENCE®  
MOSCONC CENTER · SAN FRANCISCO, CA  
MARCH 2-6, 2015 · EXPO: MARCH 4-6, 2015



# Evolution of our PS4 engine

- Solve the issues and limitations with our job system
  - All code need to be able to be jobified
  - Fibers
- CPU utilization is nowhere near 100%!
  - Process multiple frames at once
  - Memory lifetime



# Background: PS3 Engine

- Single threaded engine (30Hz)
  - Game logic followed by command buffer setup
- SPUs were used as worker threads
  - Most engine systems were running on the SPUs
- Very little gameplay code ran on SPUs



# Issues with our PS3 job system

- Jobs always ran to completion without ever yielding.
  - Complex to move gameplay onto SPUs
- User of the job system had to allocate/free resources
  - Job definitions and job lists (lifetime issues)
- State of a job list was confusing
  - Possible to add jobs while the list was running/stopped
- Job synchronization through marker index in job array
  - Had to reset job array between uses because the index would get reused



# Design goals for new job system

- Allow jobifying code that couldn't be moved to SPUs
- Jobs can yield to other jobs in the middle of execution
  - Example: Player update with kick and wait for ray casts
- Easy to use API for gameplay programmers
- No memory management for the user
- One simple way to synchronize/chain jobs
- Performance was secondary to ease-of-use of the API



# Fibers

- Like a partial thread
  - User provided stack space
  - Small context containing state of fiber and saved registers
- Executed by a thread
- Cooperative multi-threading (no preemption)
  - Switching between fibers is explicit (sceFiberSwitch on PS4)
    - Other operating systems have similar functionality
- Minimal overhead
  - No thread context switching when changing between fibers. Only register save/restore. (program counter, stack pointer, gprs...)

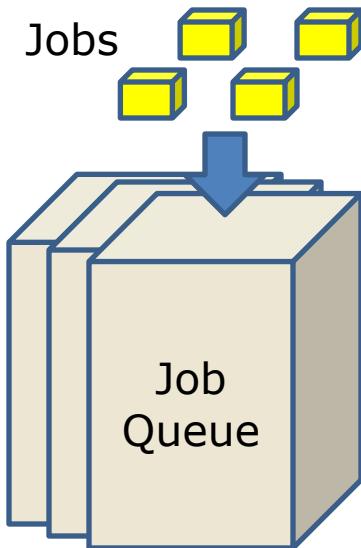


# Our job system

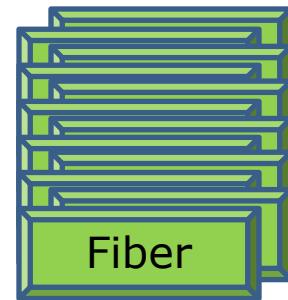
- 6 worker threads
  - Each one is locked to a CPU core
- A thread is the execution unit, the fiber is the context.
- A job always executes within the context of a fiber
- Atomic counters used for synchronization
- Fibers
  - 160 fibers (128 x 64 KiB stack, 32 x 512 KiB stack)
- 3 job queues (Low/Normal/High priority)
  - No job stealing



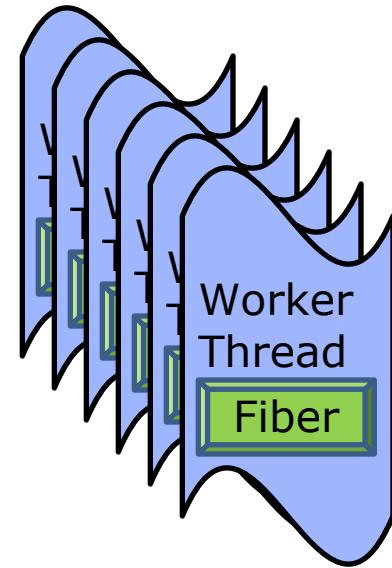
# Job System



3 Job Queues  
Low, Normal, High  
Priority



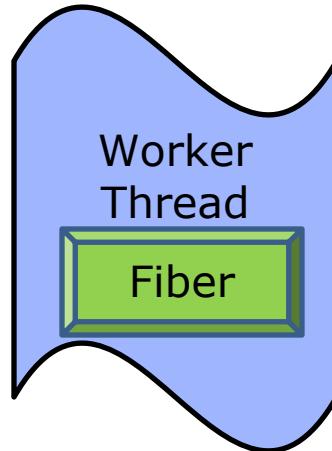
Fiber Pool - 160 Fibers  
Stack & Registers



6 Worker Threads

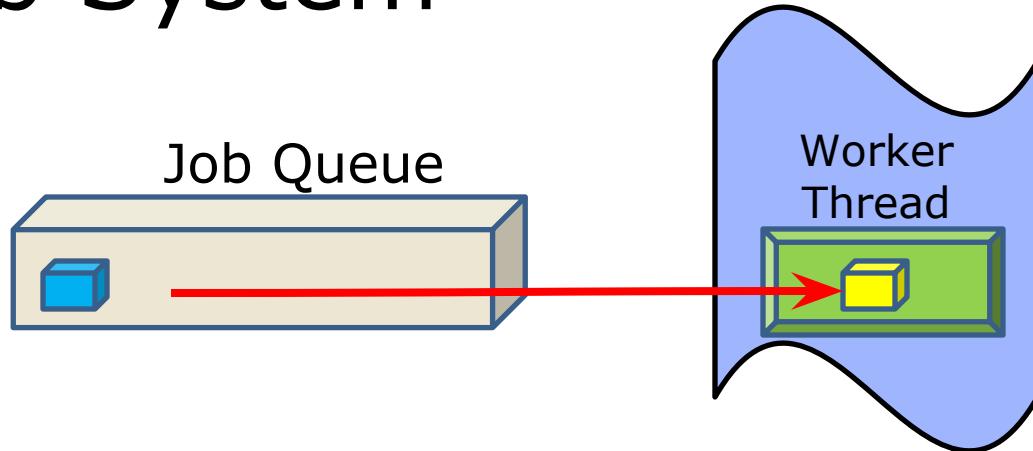


# Job System





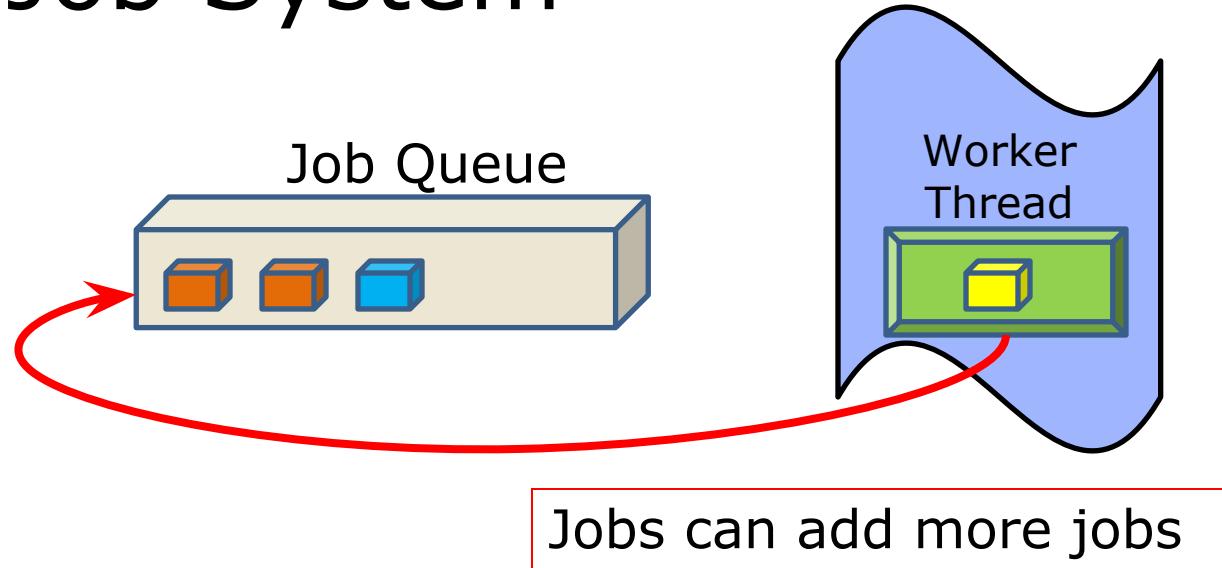
# Job System



Jobs always execute  
inside a fiber context

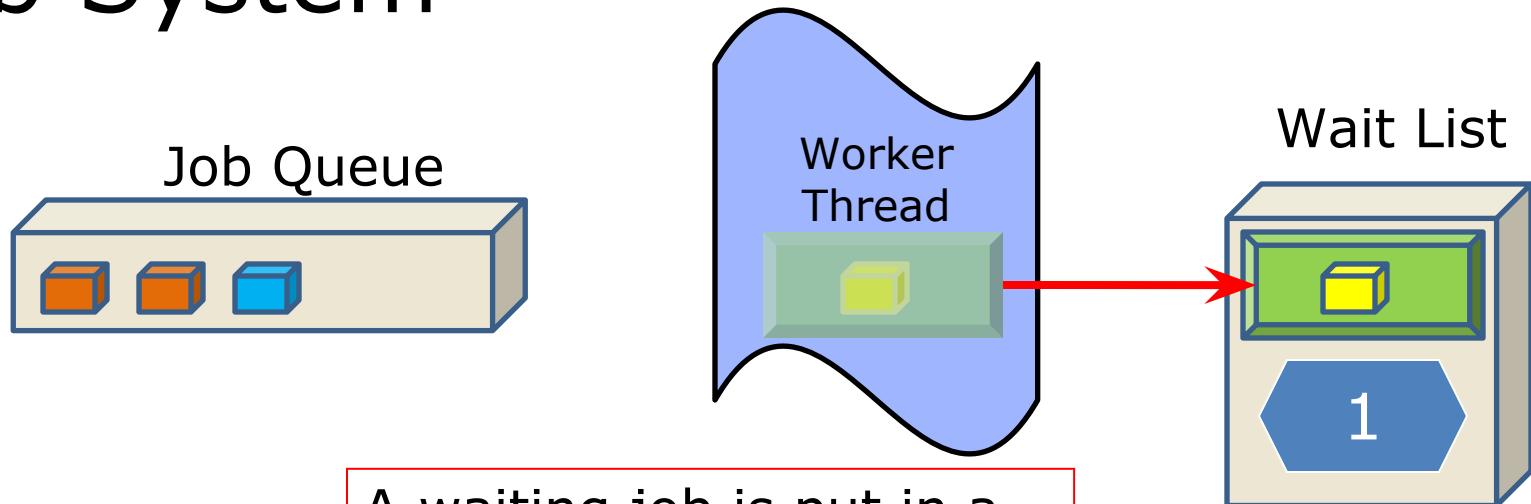


# Job System





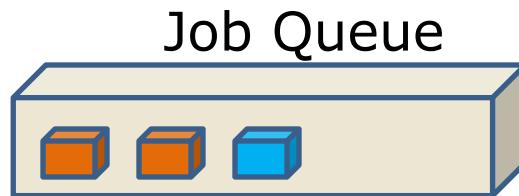
# Job System



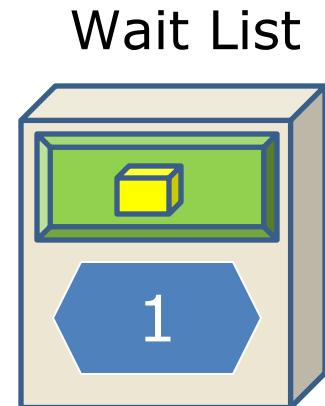
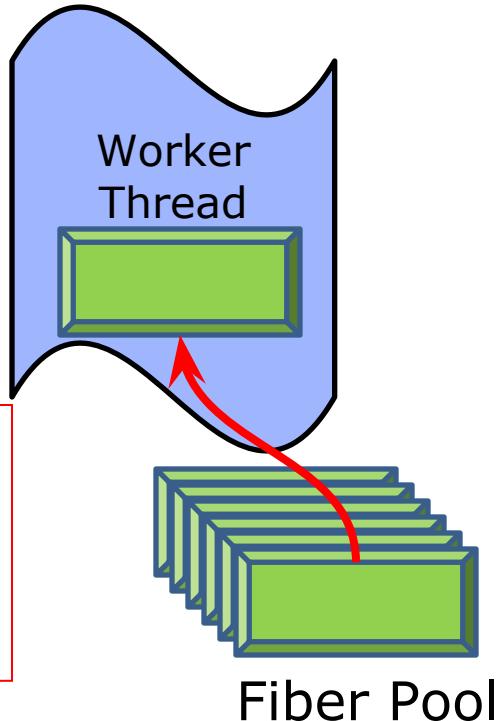
A waiting job is put in a wait list associated with the counter waited on



# Job System

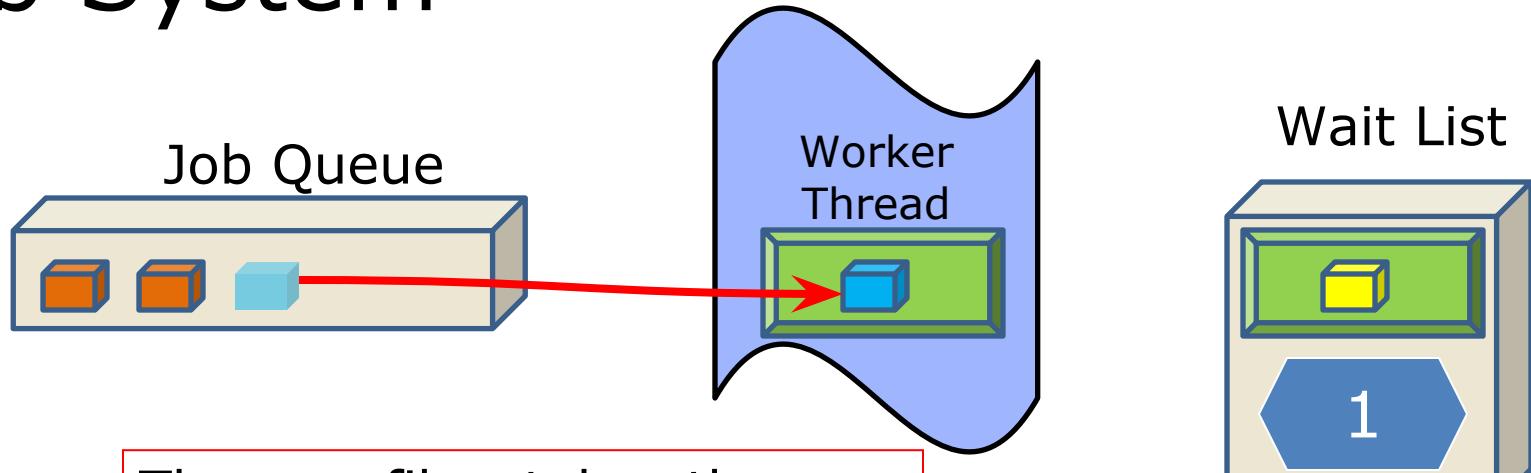


A new fiber is used to allow another job to begin executing. The stack of the waiting job is preserved in the fiber



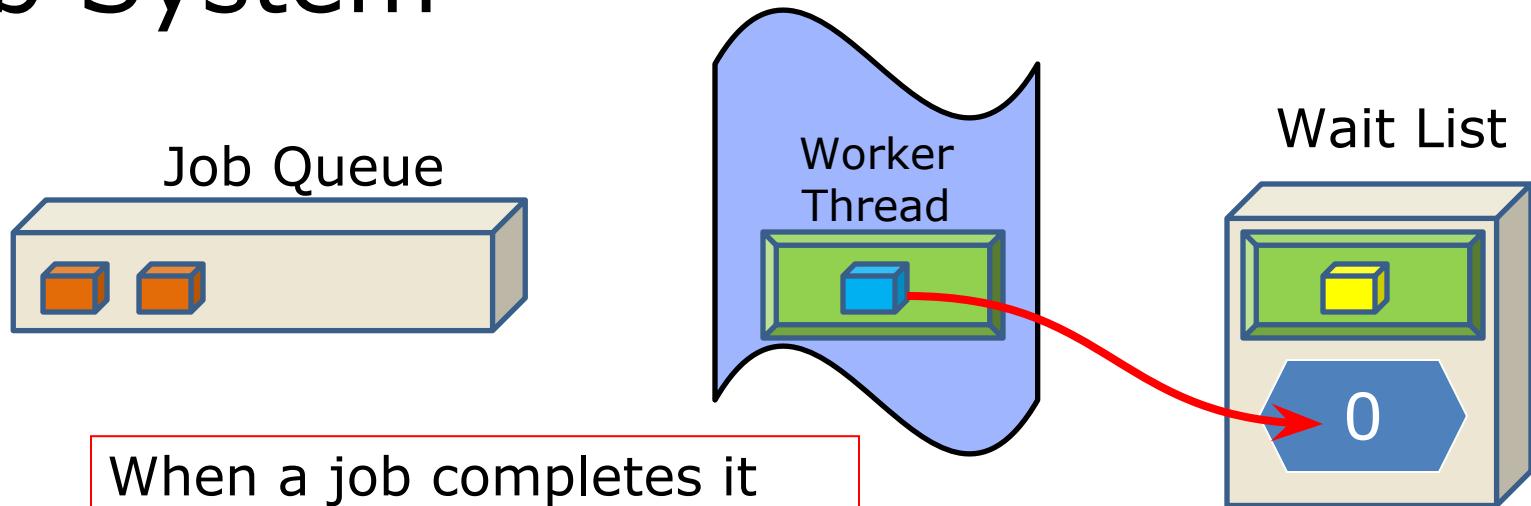


# Job System



The new fiber takes the next job from the job queue and begins executing it

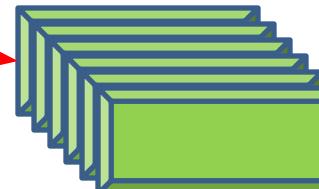
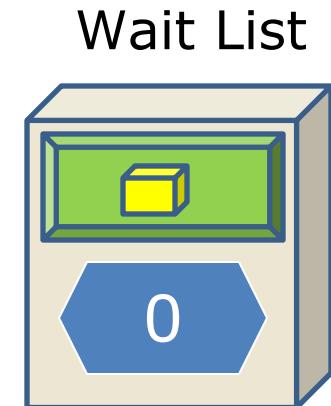
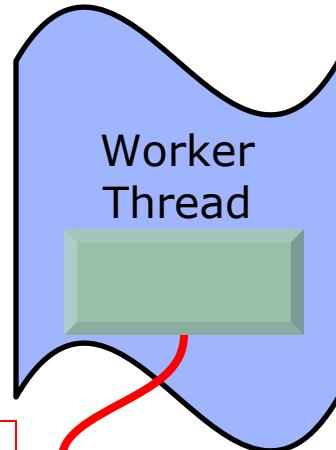
# Job System



When a job completes it will decrement an associated counter and wake up any waiting jobs



# Job System

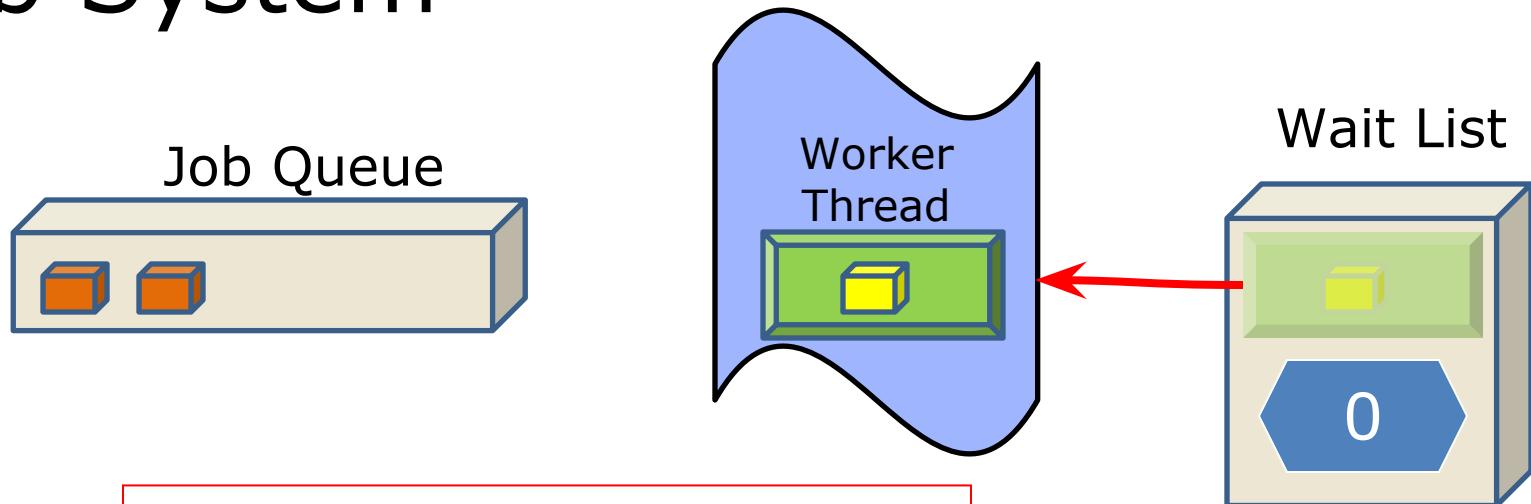


Fiber Pool

When a job is resumed we first free the current fiber



# Job System



We then switch to the waiting fiber and resume execution



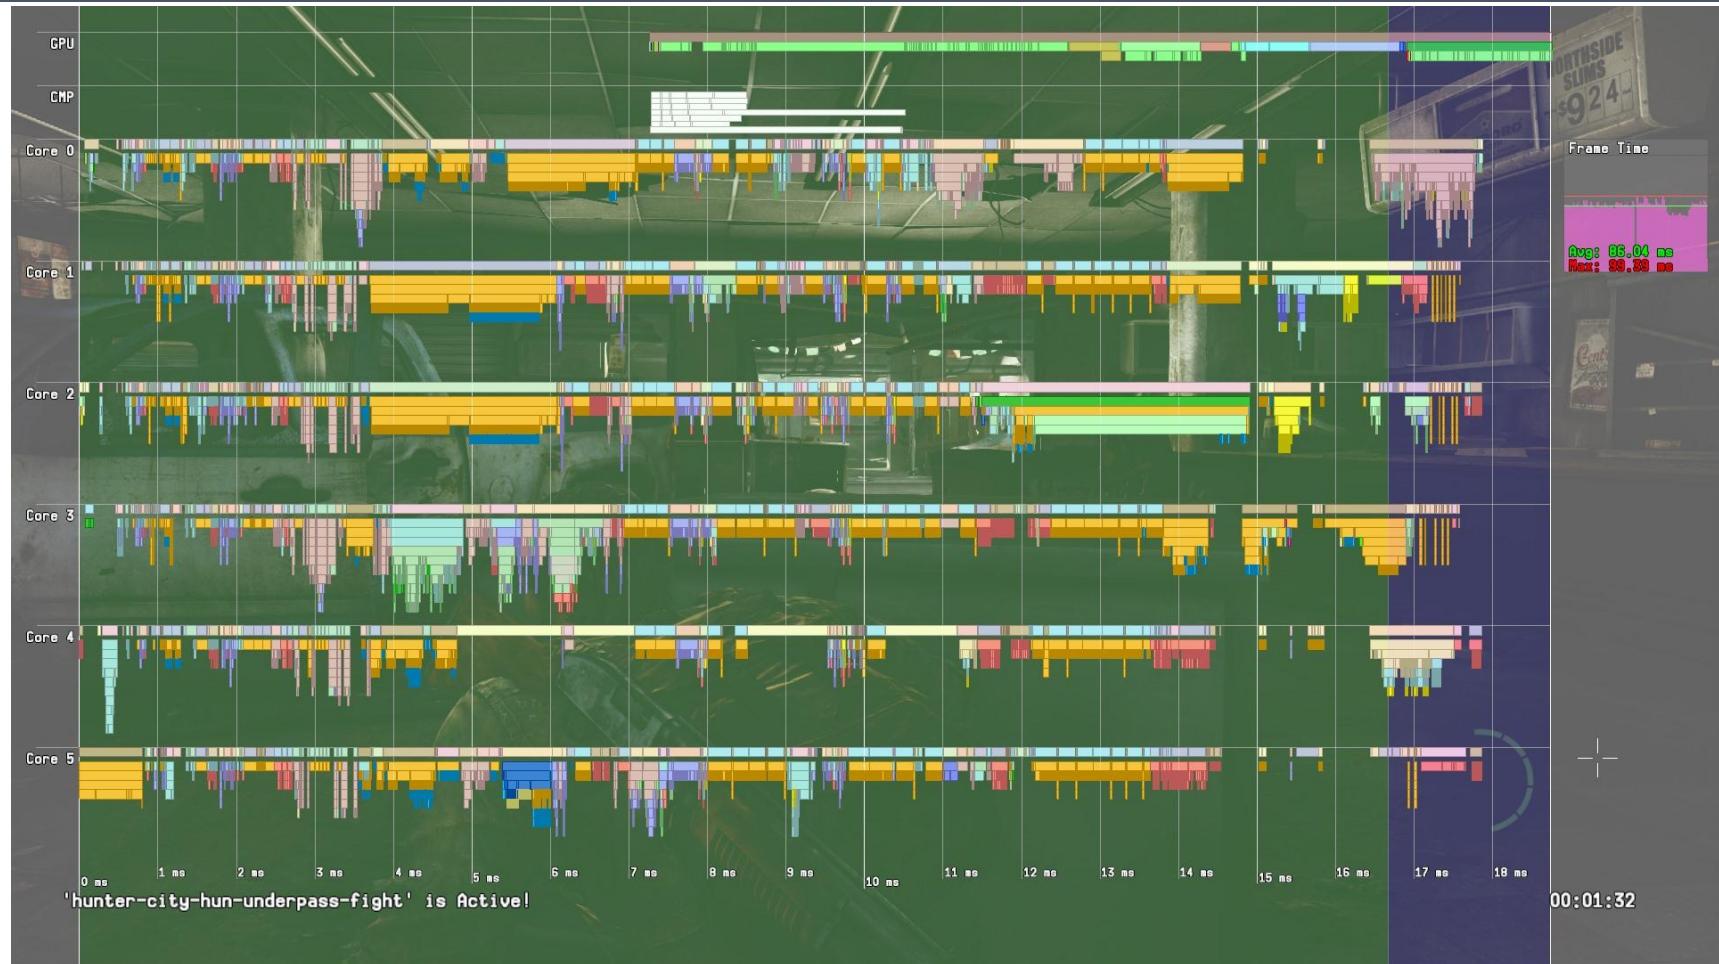
# Job System Details

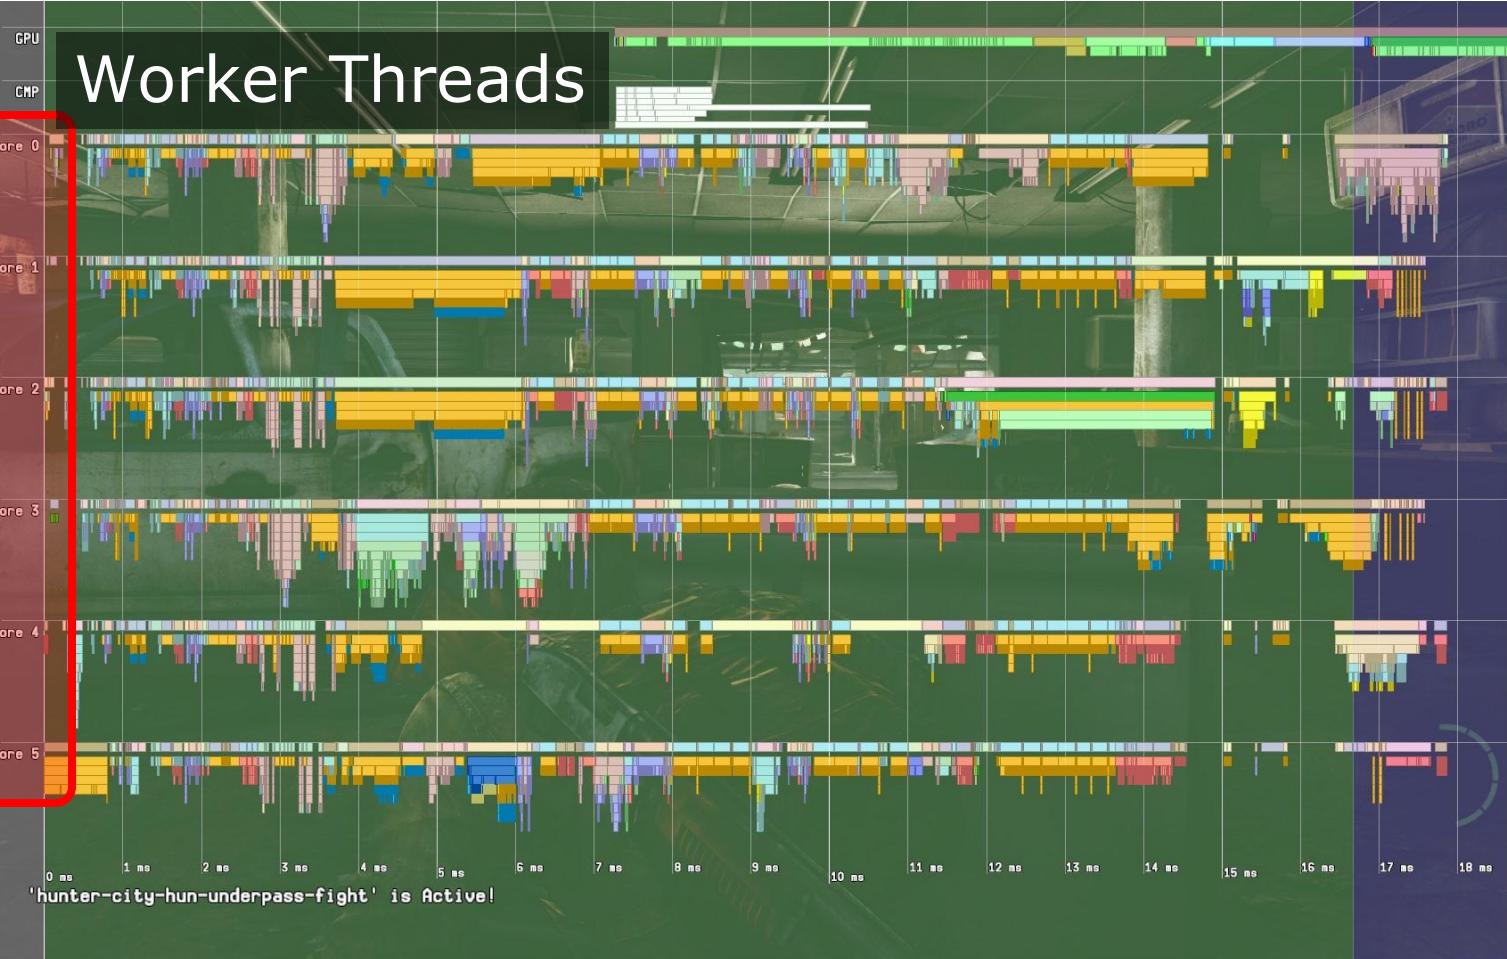
- Worker threads are locked to cores
  - Avoid context switches and unwanted core switching
  - Kernel threads can otherwise cause ripple effects across the cores
- `ndjob::WaitForCounter(counter, 0)`
  - Can be waited on by any job (moves fiber to wait list)
  - Only way of synchronizing jobs
  - Being able to yield a job at any time is extremely powerful!
- ~800-1000 jobs per frame in 'The Last of Us: Remastered'

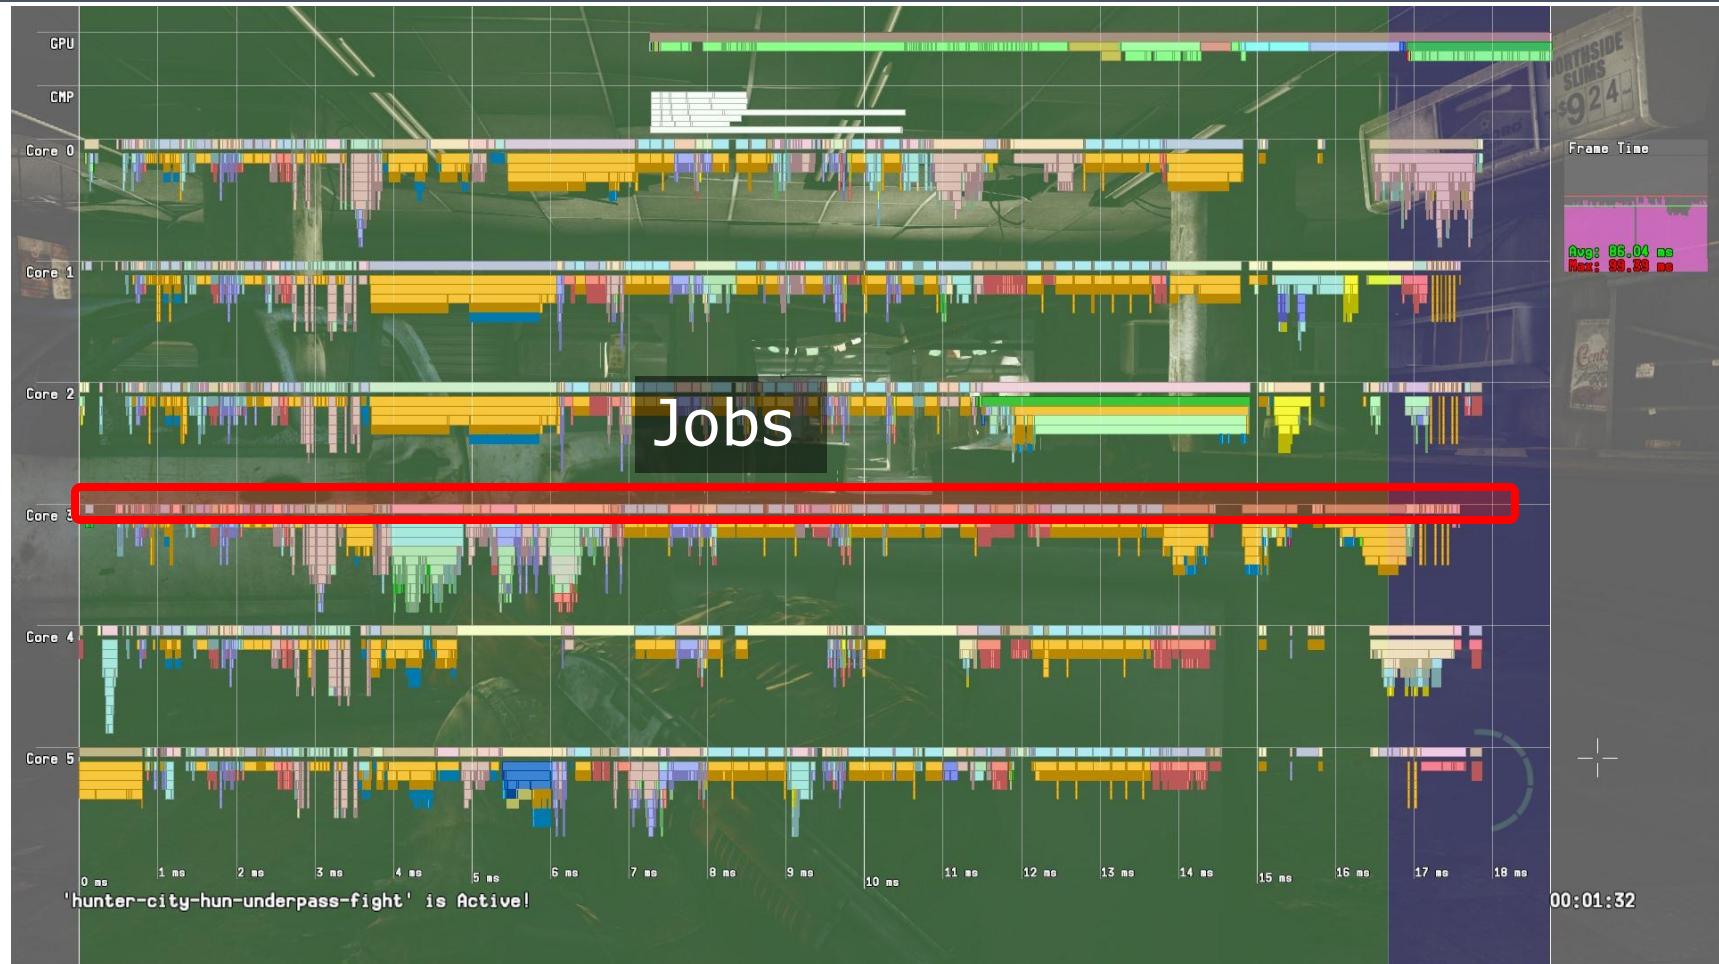


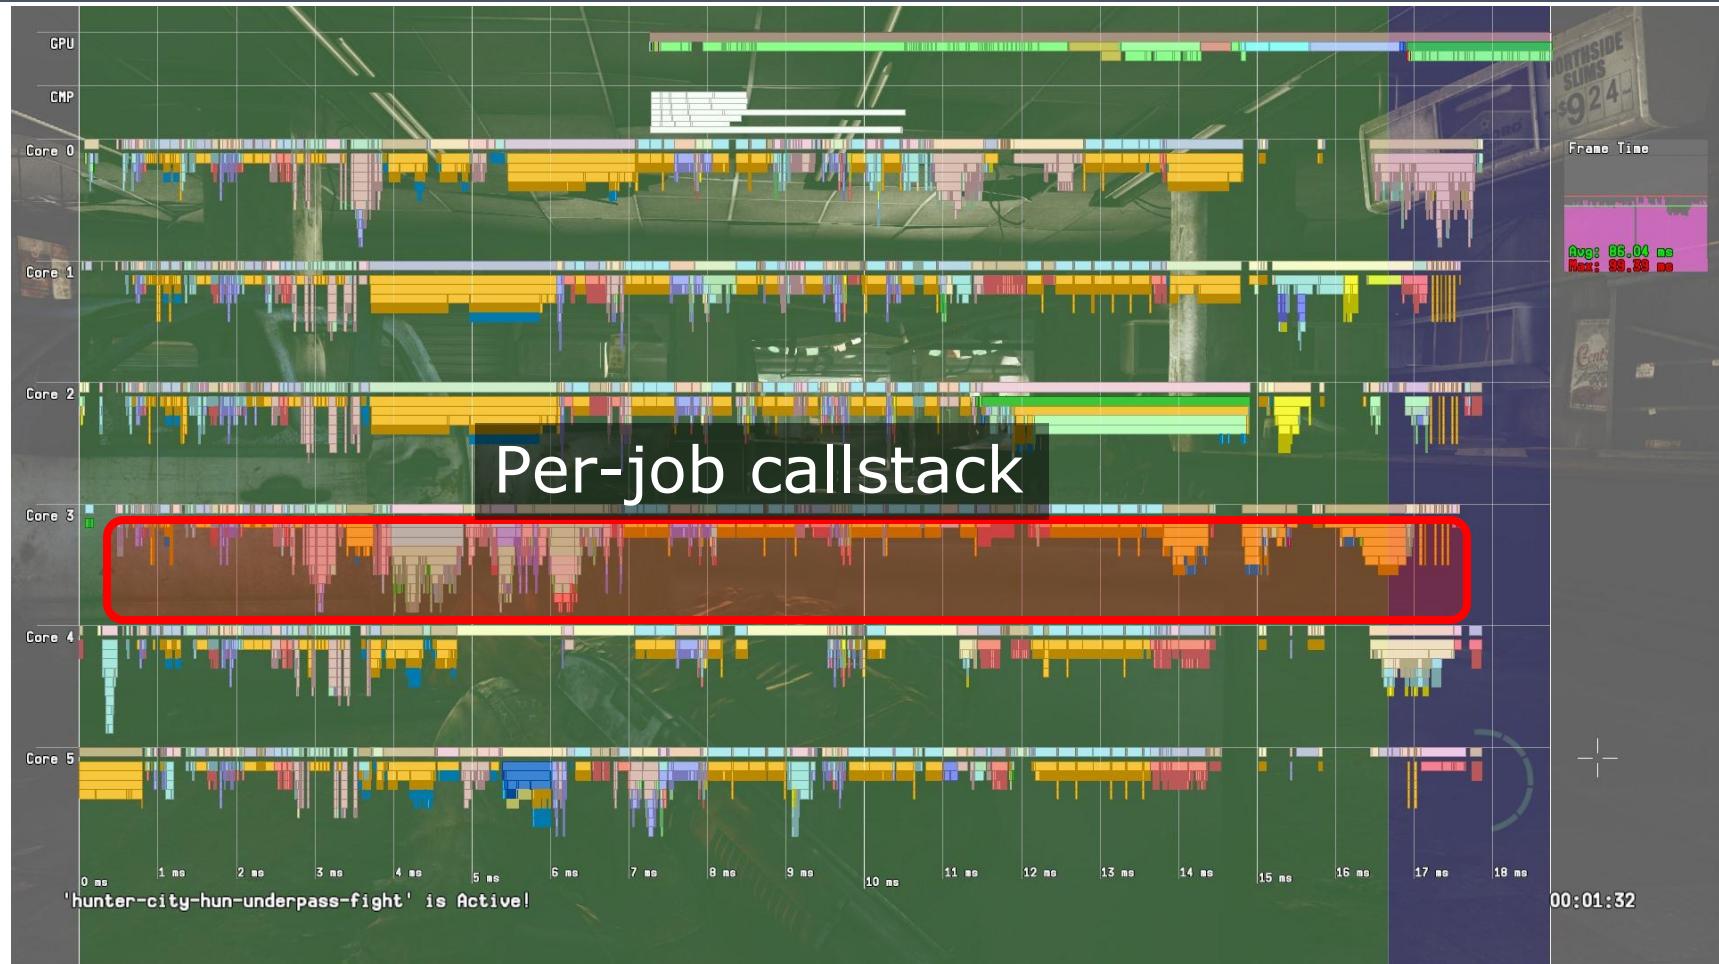
# Everything is a job

- Game object updates
- Animation updates & joint blending
- Ray casts
- Command buffer generation
- Except "I/O threads" (sockets, file I/O, system calls...)
  - These are system threads
  - Implemented like interrupt handlers (Read data, post new job)
  - Always waiting and never do expensive processing of the data











# Example – Animate All Objects

```
Object g_Objects[100];

void AnimateObject(void* pObj)
{
    ...
}

void AnimateAllObjects()
{
    for (int objIndex = 0; objIndex < 100; ++objIndex)
    {
        AnimateObject(&g_Objects[objIndex]);
    }
}
```



# Example – Animate All Objects

```
Object g_Objects[100];  
  
void AnimateObject(void* pObj)  
{  
    ...  
}  
  
void AnimateAllObjects()  
{  
    for (int objIndex = 0; objIndex < 100; ++ objIndex )  
    {  
        AnimateObject(&g_Objects[objIndex]);  
    }  
}
```



# Example – Animate All Objects

```
Object g_Objects[100];

JOB_ENTRY_POINT(AnimateObjectJob)
{
    ...
}

void AnimateAllObjects()
{
    JobDecl jobDecls[100];
    for (int jobIndex = 0; jobIndex < 100; ++jobIndex)
    {
        jobDecls[jobIndex] = JobDecl(AnimateObjectJob, &g_Objects[jobIndex]);
    }
    Counter* pJobCounter = NULL;
    RunJobs(&jobDecls, 100, &pJobCounter);
    WaitForCounterAndFree(pJobCounter, 0);
}
```



# Example – Animate All Objects

```
Object g_Objects[100];  
  
JOB_ENTRY_POINT(AnimateObjectJob)  
{  
    ...  
}  
  
void AnimateAllObjects()  
{  
    JobDecl jobDecls[100];  
    for (int jobIndex = 0; jobIndex < 100; ++jobIndex)  
    {  
        jobDecls[jobIndex] = JobDecl(AnimateObjectJob, &g_Objects[jobIndex]);  
    }  
    Counter* pJobCounter = NULL;  
    RunJobs(&jobDecls, 100, &pJobCounter);  
    WaitForCounterAndFree(pJobCounter, 0);  
}
```

Jobs can be created on  
the stack



# Example – Animate All Objects

```
Object g_Objects[100];  
  
JOB_ENTRY_POINT(AnimateObjectJob)  
{  
    ...  
}  
  
void AnimateAllObjects()  
{  
    JobDecl jobDecls[100];  
    for (int jobIndex = 0; jobIndex < 100; ++jobIndex)  
    {  
        jobDecls[jobIndex] = JobDecl(AnimateObjectJob, &g_Objects[jobIndex]);  
    }  
    Counter* pJobCounter = NULL;  
    RunJobs(&jobDecls, 100, &pJobCounter);  
    WaitForCounterAndFree(pJobCounter, 0);  
}
```

Fill in the job declarations  
with entry point and  
parameter



# Example – Animate All Objects

```
Object g_Objects[100];
```

```
JOB_ENTRY_POINT(AnimateObjectJob)
{
    ...
}
```

```
void AnimateAllObjects()
{
    JobDecl jobDecls[100];
    for (int jobIndex = 0; jobIndex < 100; ++jobIndex)
    {
        jobDecls[jobIndex] = JobDecl(AnimateObjectJob, &g_Objects[jobIndex]);
    }
    Counter* pJobCounter = NULL;
    RunJobs(&jobDecls, 100, &pJobCounter);
    WaitForCounterAndFree(pJobCounter, 0);
}
```

The job entry point is  
easily defined



# Example – Animate All Objects

```
Object g_Objects[100];  
  
JOB_ENTRY_POINT(AnimateObjectJob)  
{  
    ...  
}  
  
void AnimateAllObjects()  
{  
    JobDecl jobDecls[100];  
    for (int jobIndex = 0; jobIndex < 100; ++jobIndex)  
    {  
        jobDecls[jobIndex] = JobDecl(AnimateObjectJob, &g_Objects[jobIndex]);  
    }  
Counter* pJobCounter = NULL;  
RunJobs(&jobDecls, 100, &pJobCounter);  
WaitForCounterAndFree(pJobCounter, 0);  
}
```

Schedule all jobs and receive a counter that you can wait for



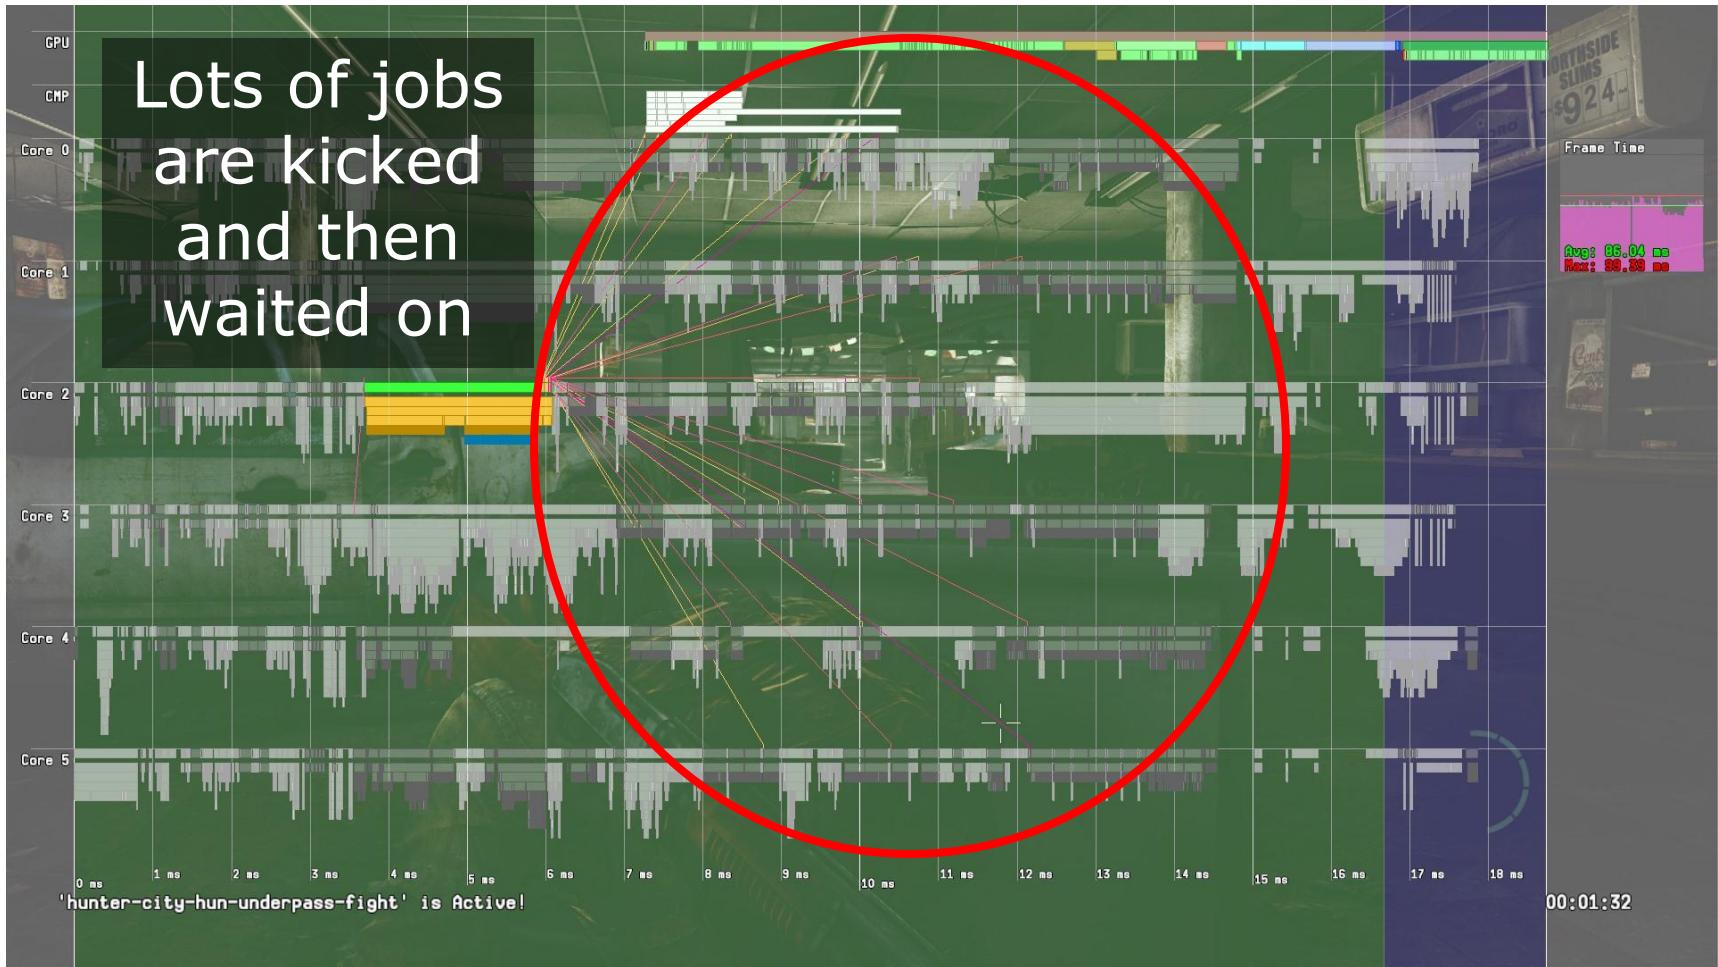
# Example – Animate All Objects

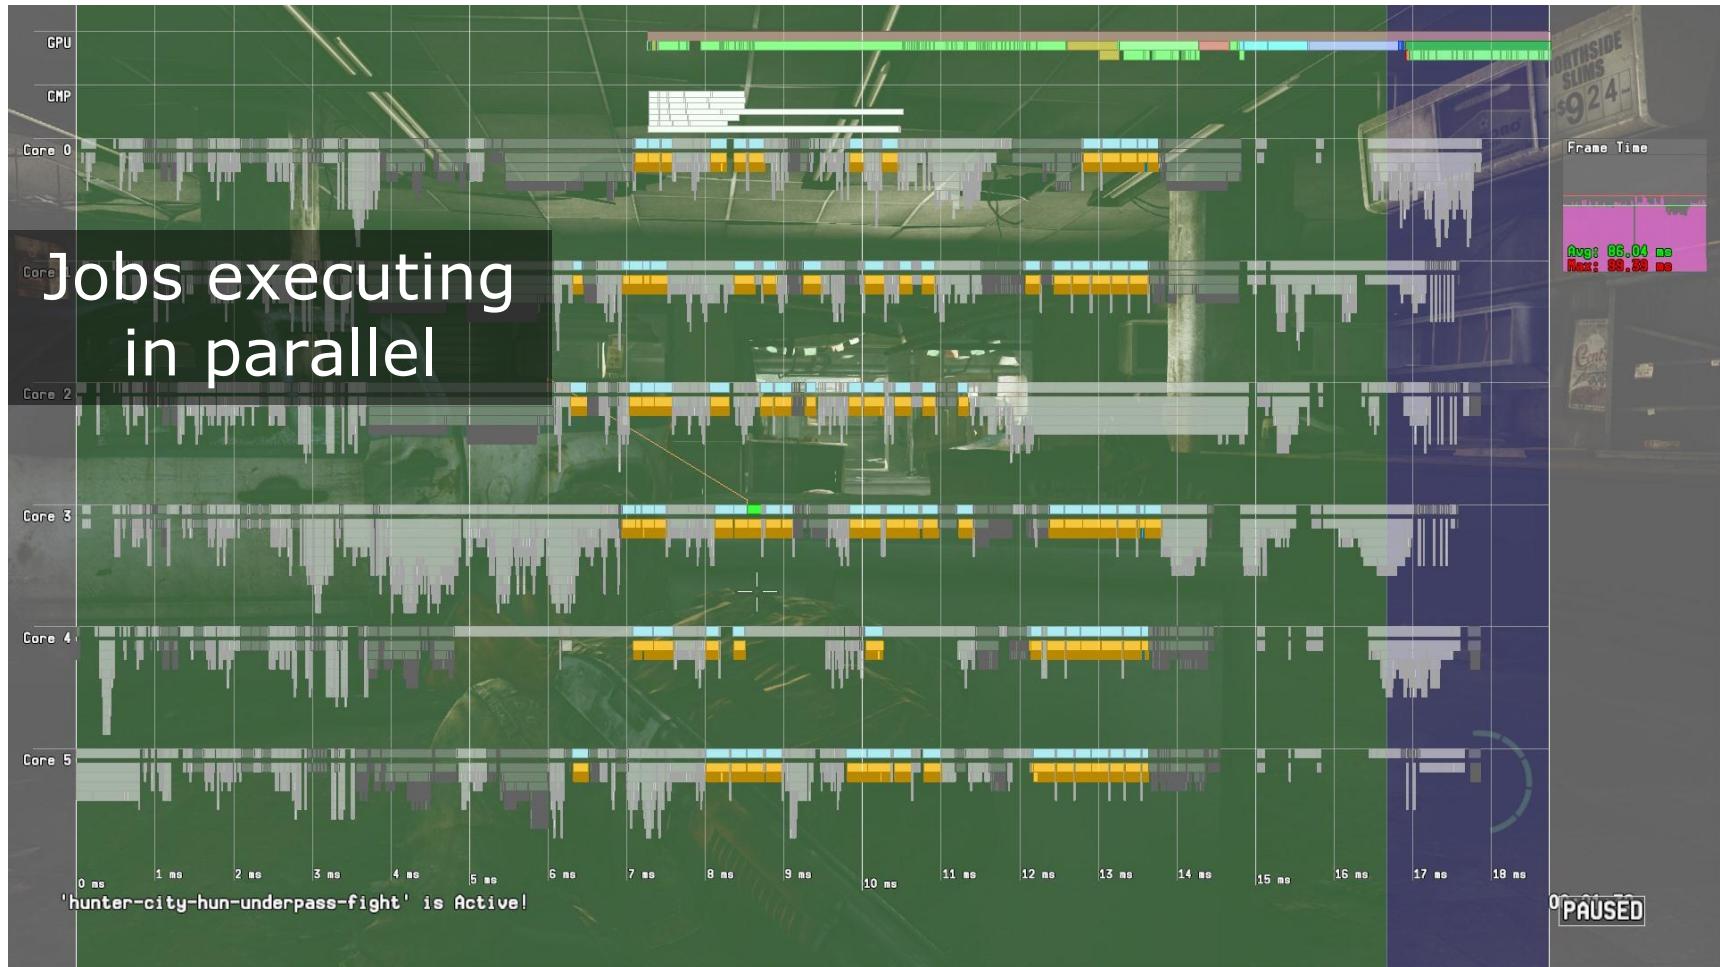
```
Object g_Objects[100];

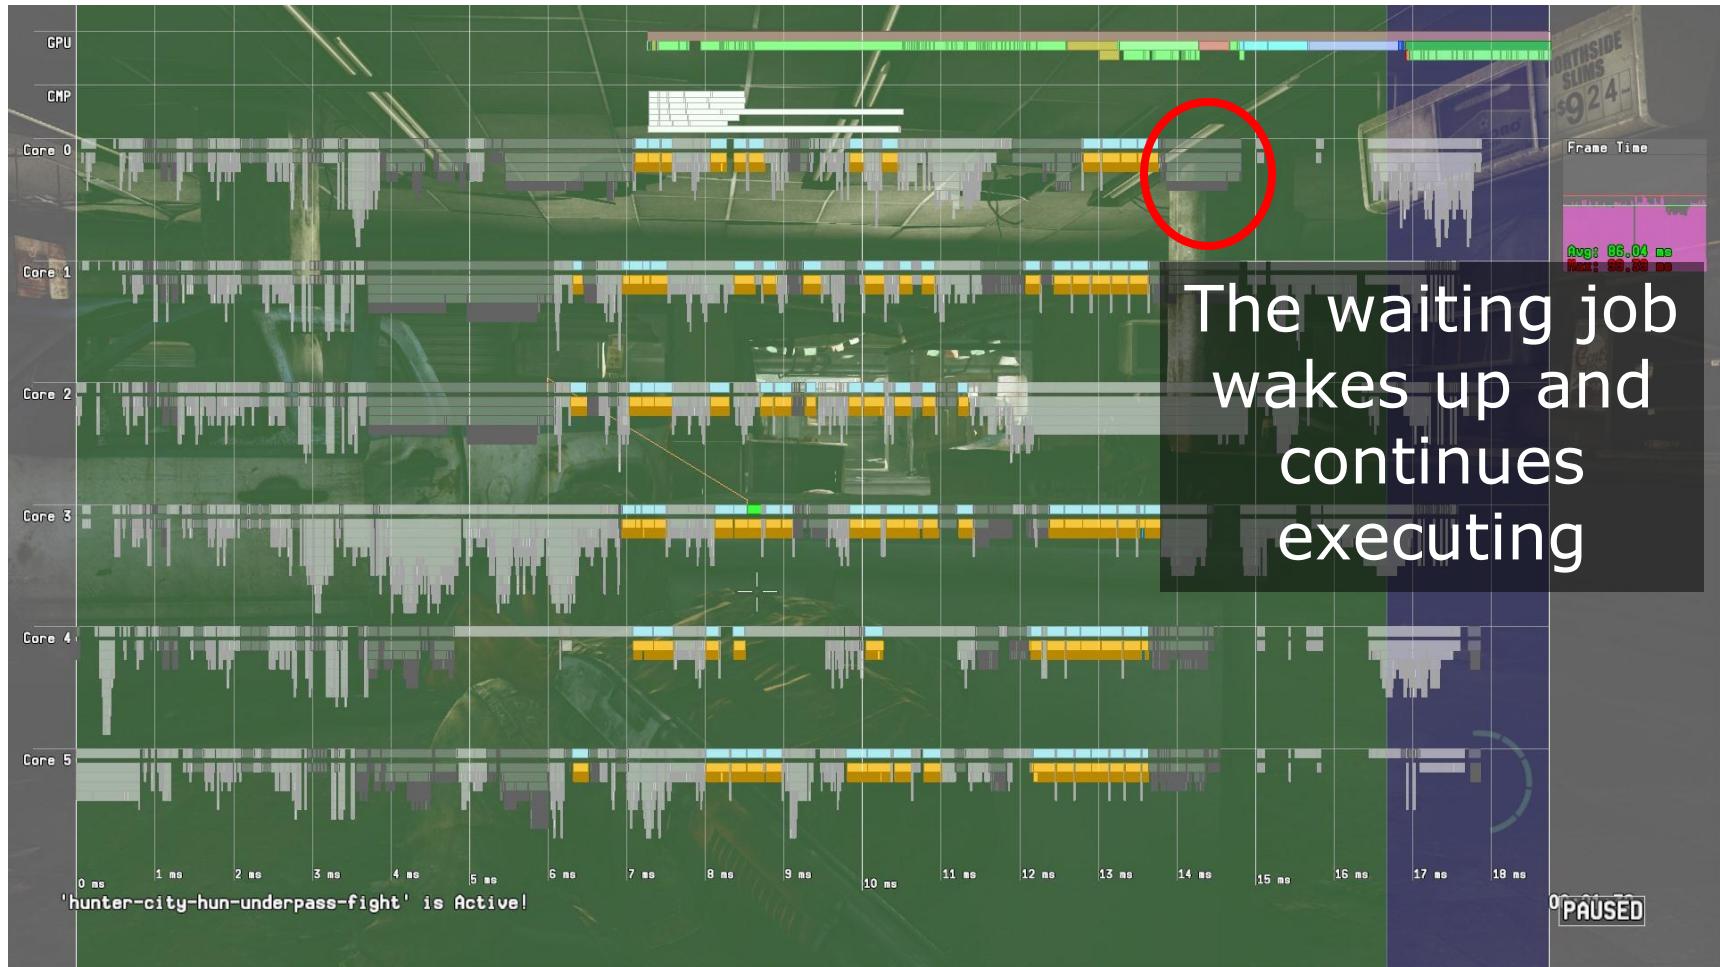
JOB_ENTRY_POINT(AnimateObjectJob)
{
    ...
}

void AnimateAllObjects()
{
    JobDecl jobDecls[100];
    for (int jobIndex = 0; jobIndex < 100; ++jobIndex)
    {
        jobDecls[jobIndex] = JobDecl(AnimateObjectJob, &g_Objects[jobIndex]);
    }
    Counter* pJobCounter = NULL;
    RunJobs(&jobDecls, 100, &pJobCounter);
WaitForCounterAndFree(pJobCounter, 0);
}
```

Wait for all jobs by waiting  
for the counter to be zero.  
Job goes to sleep.









# Pros of new job system

- Extremely easy to jobify existing gameplay updates
  - Deep call stacks are no problem
- Having one job wait for another is straight-forward
  - WaitForCounter(...)
- Super-lightweight to switch fibers
  - System supported operation - sceFiberSwitch() on PS4
  - Save the program counter and stack pointer...
    - ...and all the other registers
  - Restore the program counter and stack pointer...
    - ...and all the other registers



# Cons of new job system

- System synchronization primitives can no longer be used
  - Mutex, semaphore, condition variables...
  - Locked to a particular thread. Fibers migrate between threads
- Synchronization has to be done on the hardware level
  - Atomic spin locks are used almost everywhere
  - Special job mutex is used for locks held longer
    - Puts the current job to sleep if needed instead of spin lock



# Support For Fibers

- Fibers and their call stacks are viewable in the debugger
  - Can inspect fibers just like you would inspect threads
- Fibers can be named/renamed
  - Indicate current job
- Crash handling
  - Fiber call stacks are saved in the core dumps just like threads



# Support For Fibers...

- Fiber-Safe Thread Local Storage (TLS) Optimizations
  - Issue: TLS address is allowed to be cached for the duration of the function by default. Switch fiber in the middle of function and you wake up with wrong TLS pointer. ☹
  - Currently not supported by Clang
  - Workaround: Use separate CPP file for TLS access
- Use adaptive mutexes in job system
  - Jobs can be added from normal threads
    - Spin locks -> Dead locks
  - Spin and attempt to grab lock before doing system call
  - Solves priority inversion deadlocks
  - Can avoid most system calls due to initial spin



# Other job systems

- Intel Thread Build Blocks (TBB) is popular
- Solves dependencies between jobs and chaining but requires every job to execute until completion...
  - ... unless you want context switching
- OR allow job nesting while a job is waiting.
  - This will prevent the first job from resuming until the second job has completed.
- A fiber-based job system has none of these issues

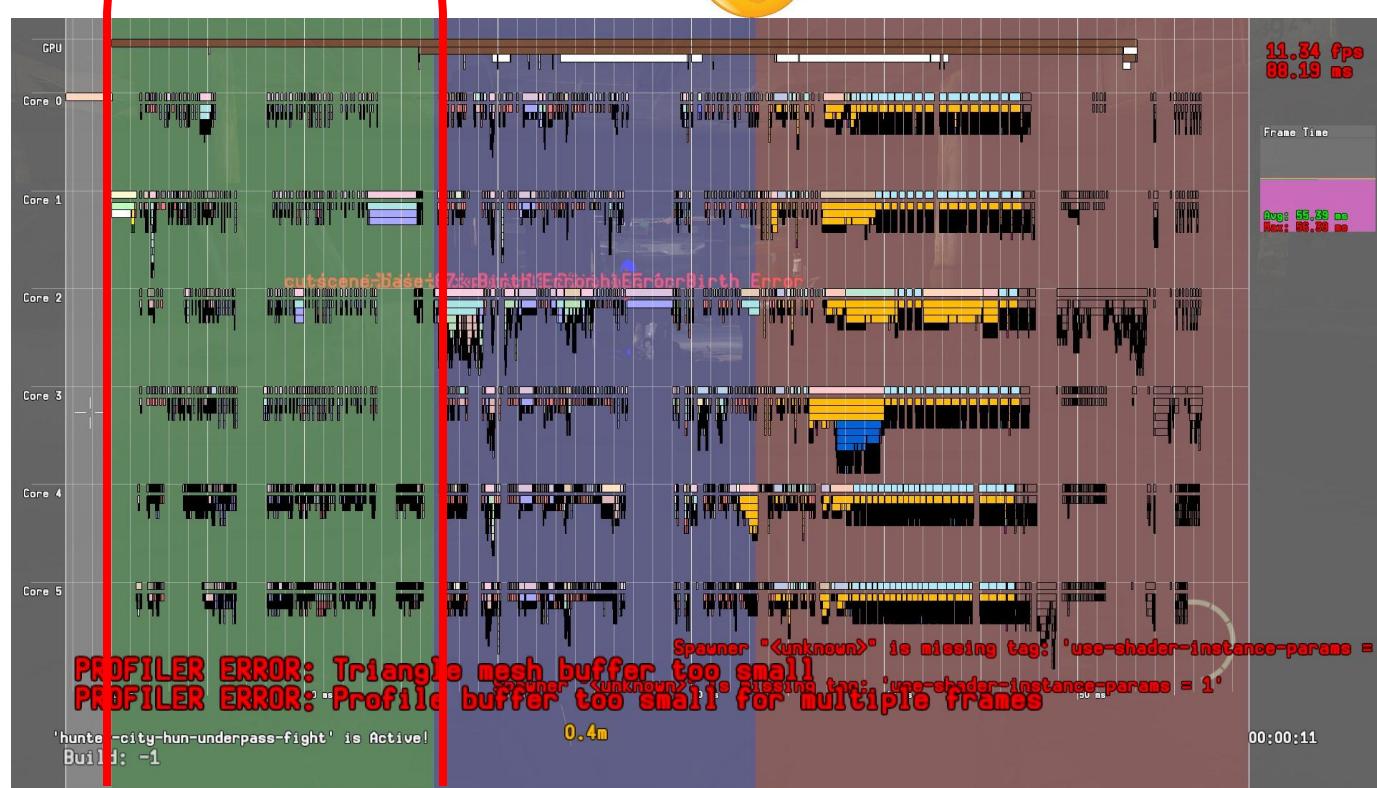


# Moving on...

- We have our new job system
- All code can now be jobified
- Let's do it!

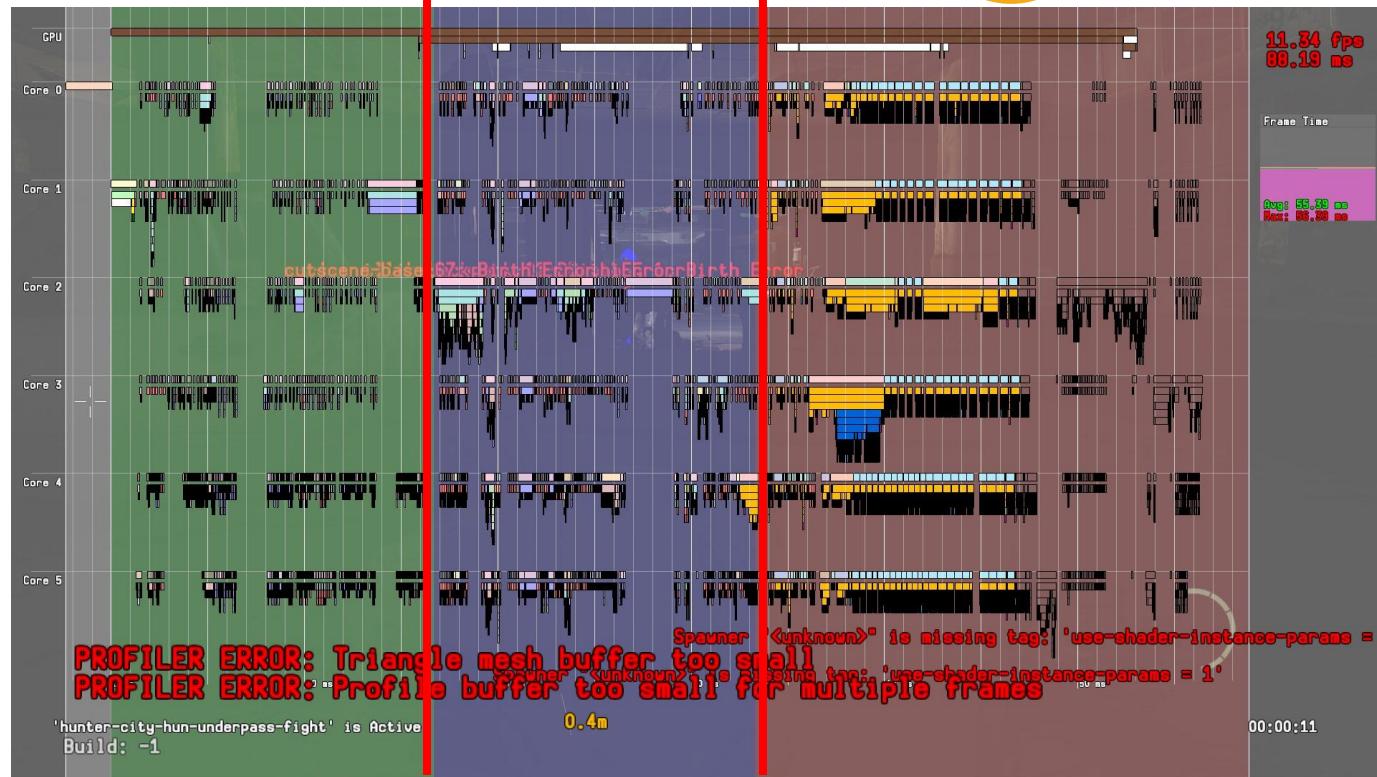


16.66 ms(60 fps)



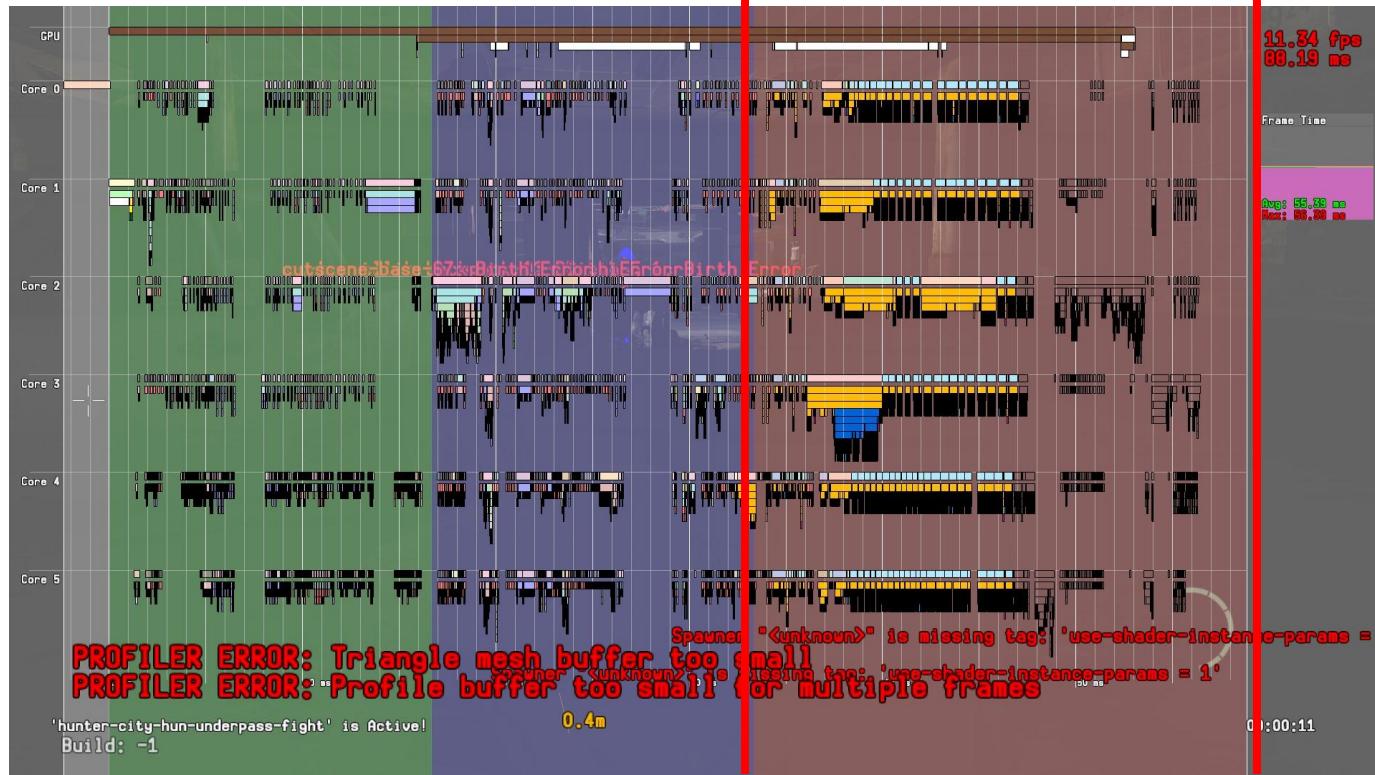


33.33 ms(30 fps)



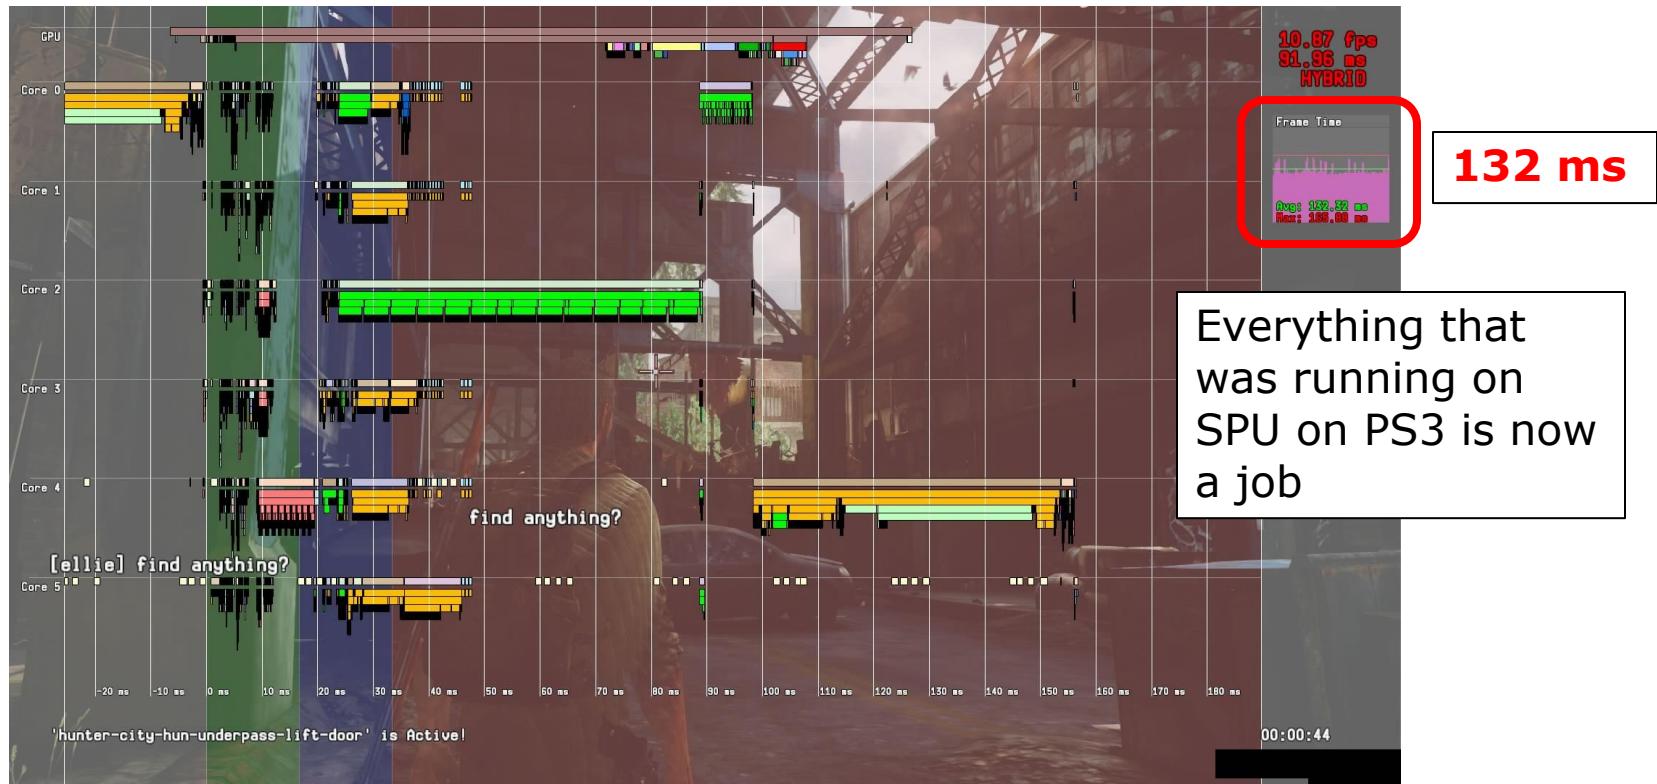


# Less than 30 fps





# Initial Port



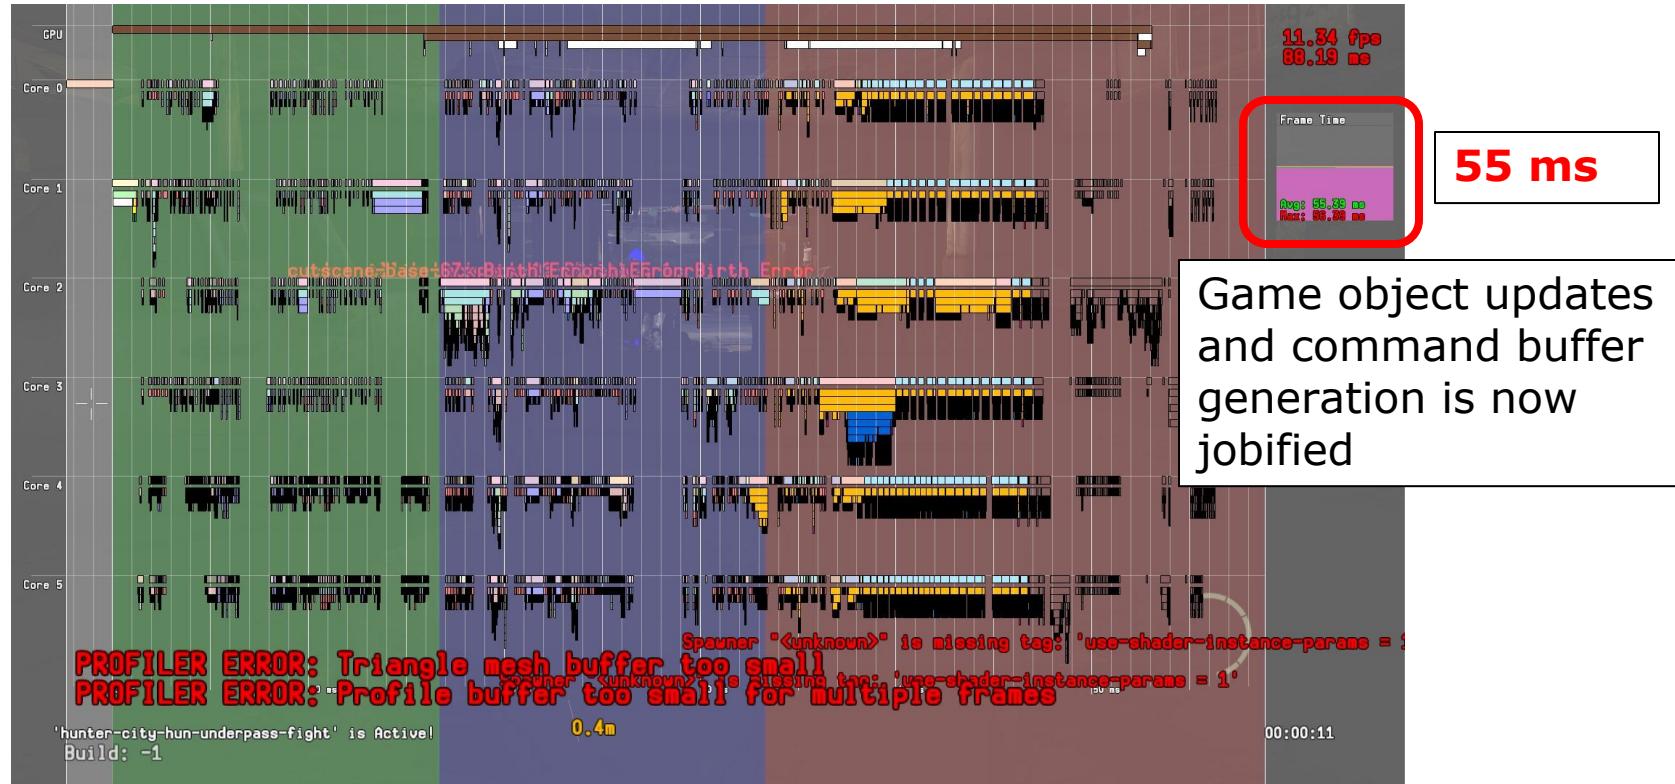


# Initial Port



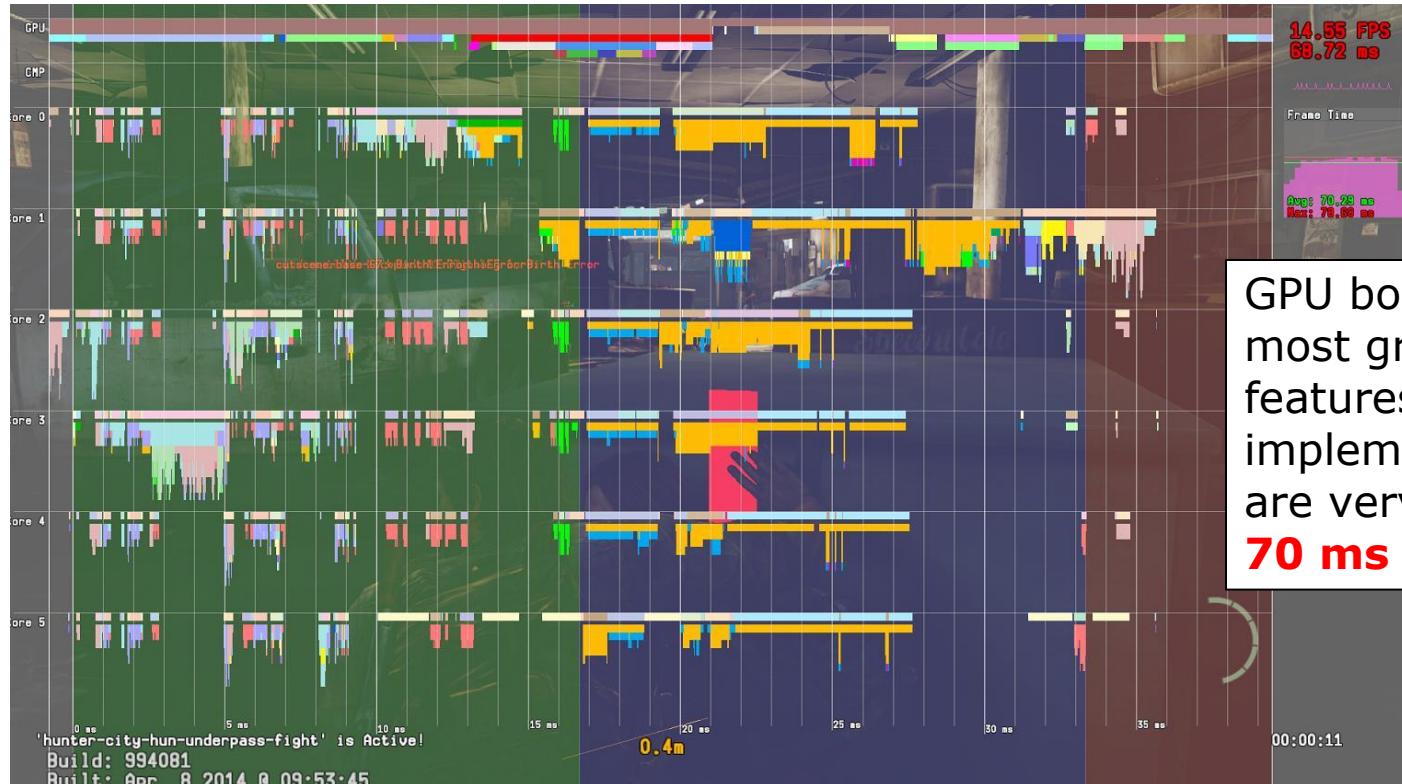


# Jobify everything





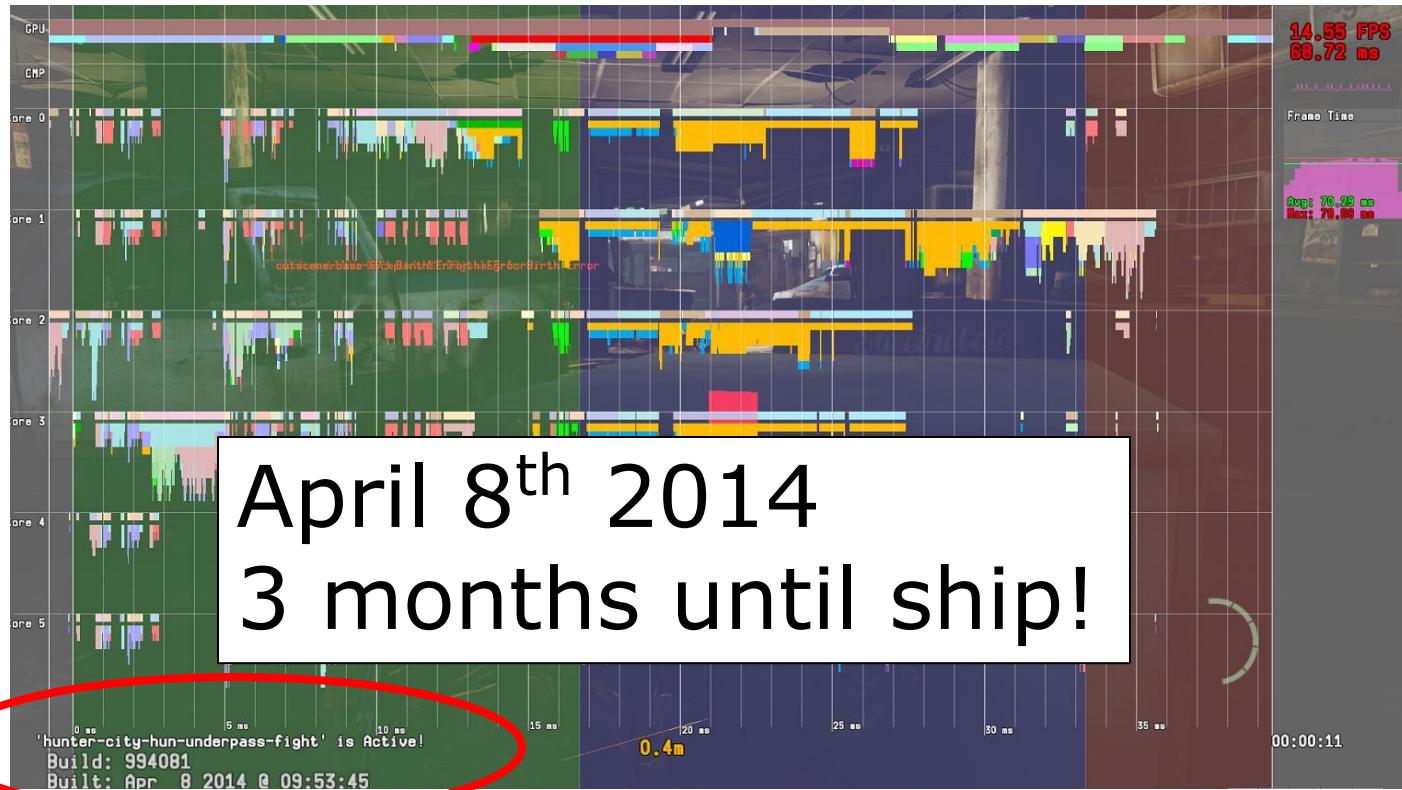
# More jobs and fewer locks



**36 ms**

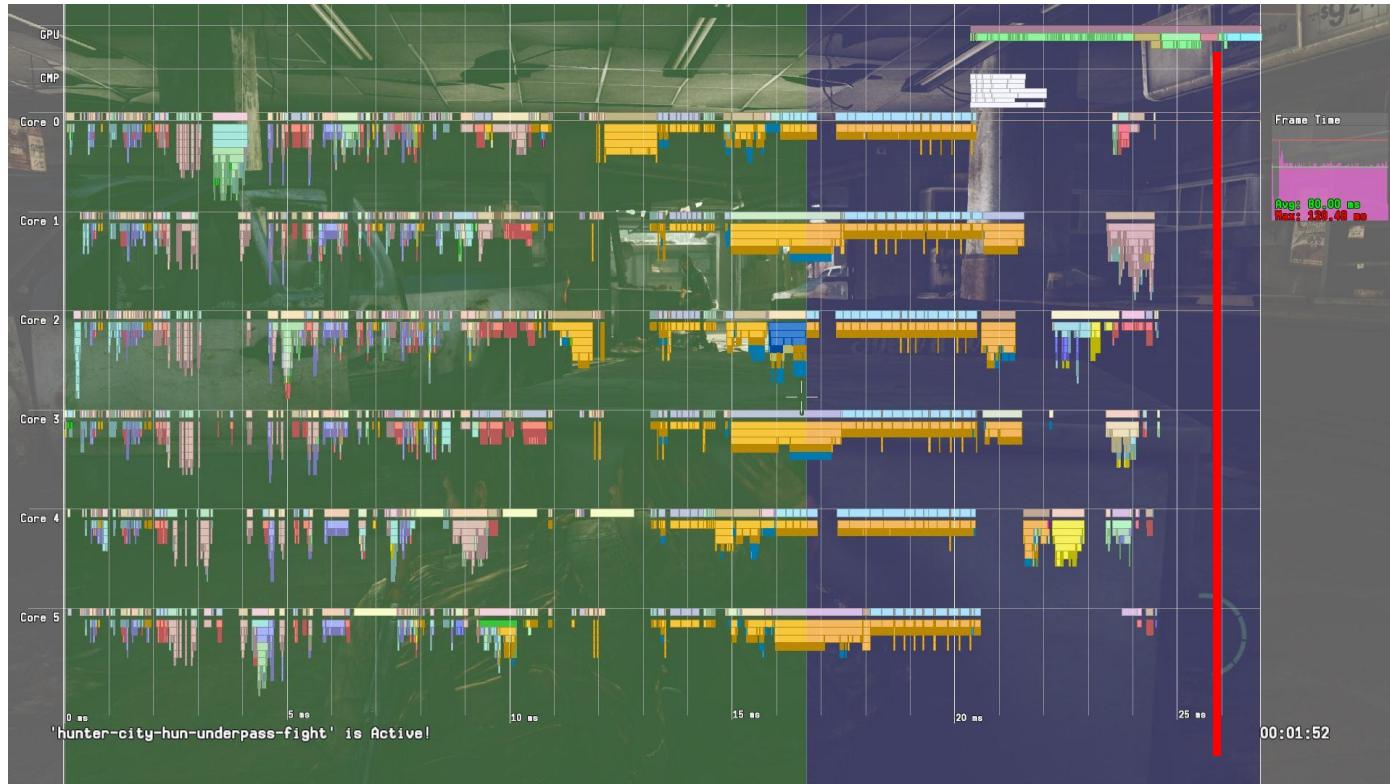
GPU bound due to  
most graphics  
features being  
implemented, but  
are very slow.

**70 ms GPU frame**





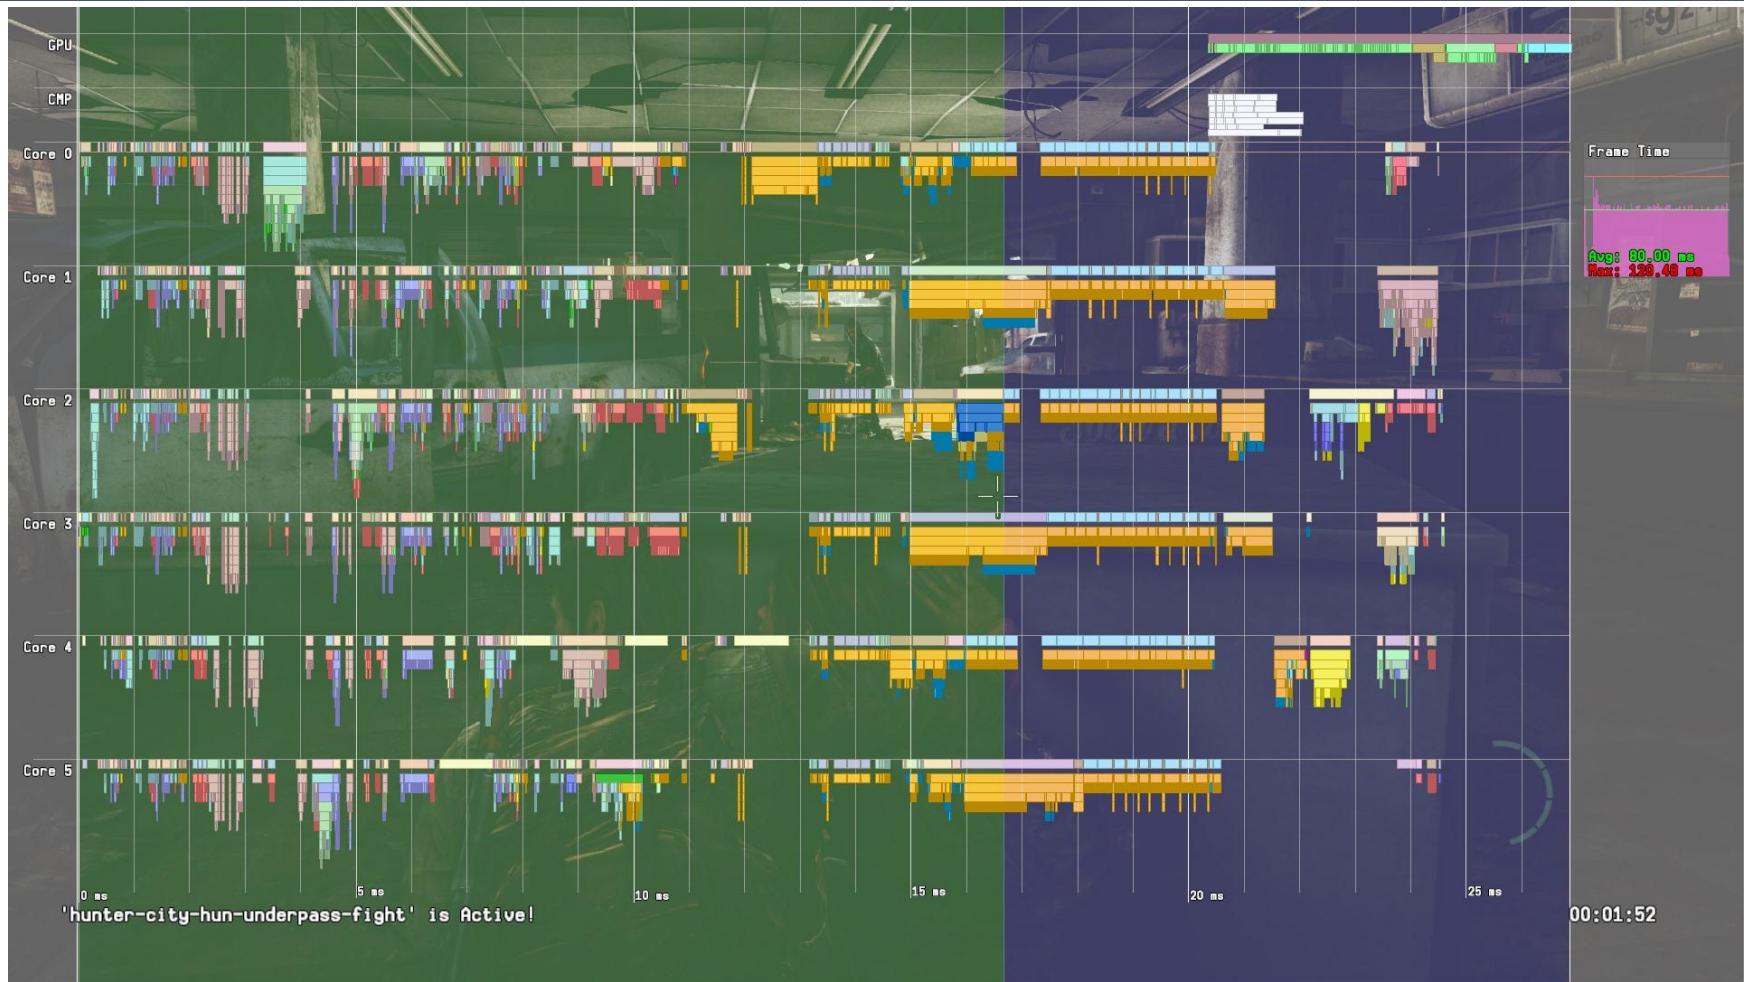
# Diminishing returns...

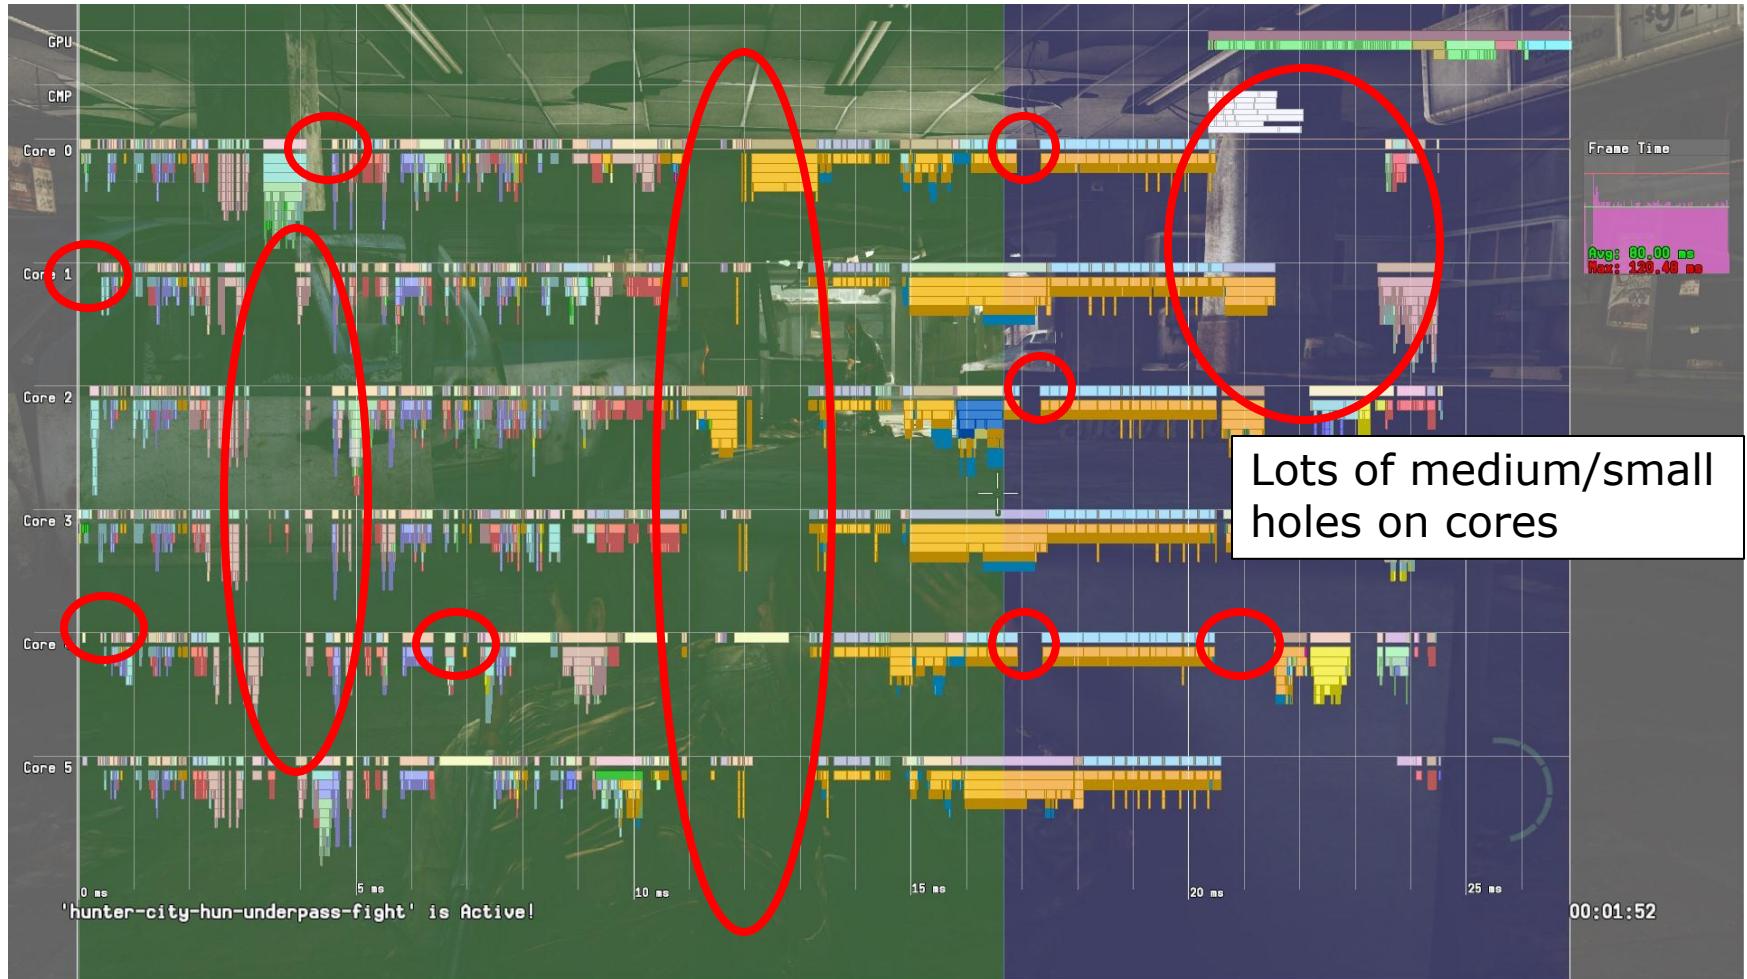


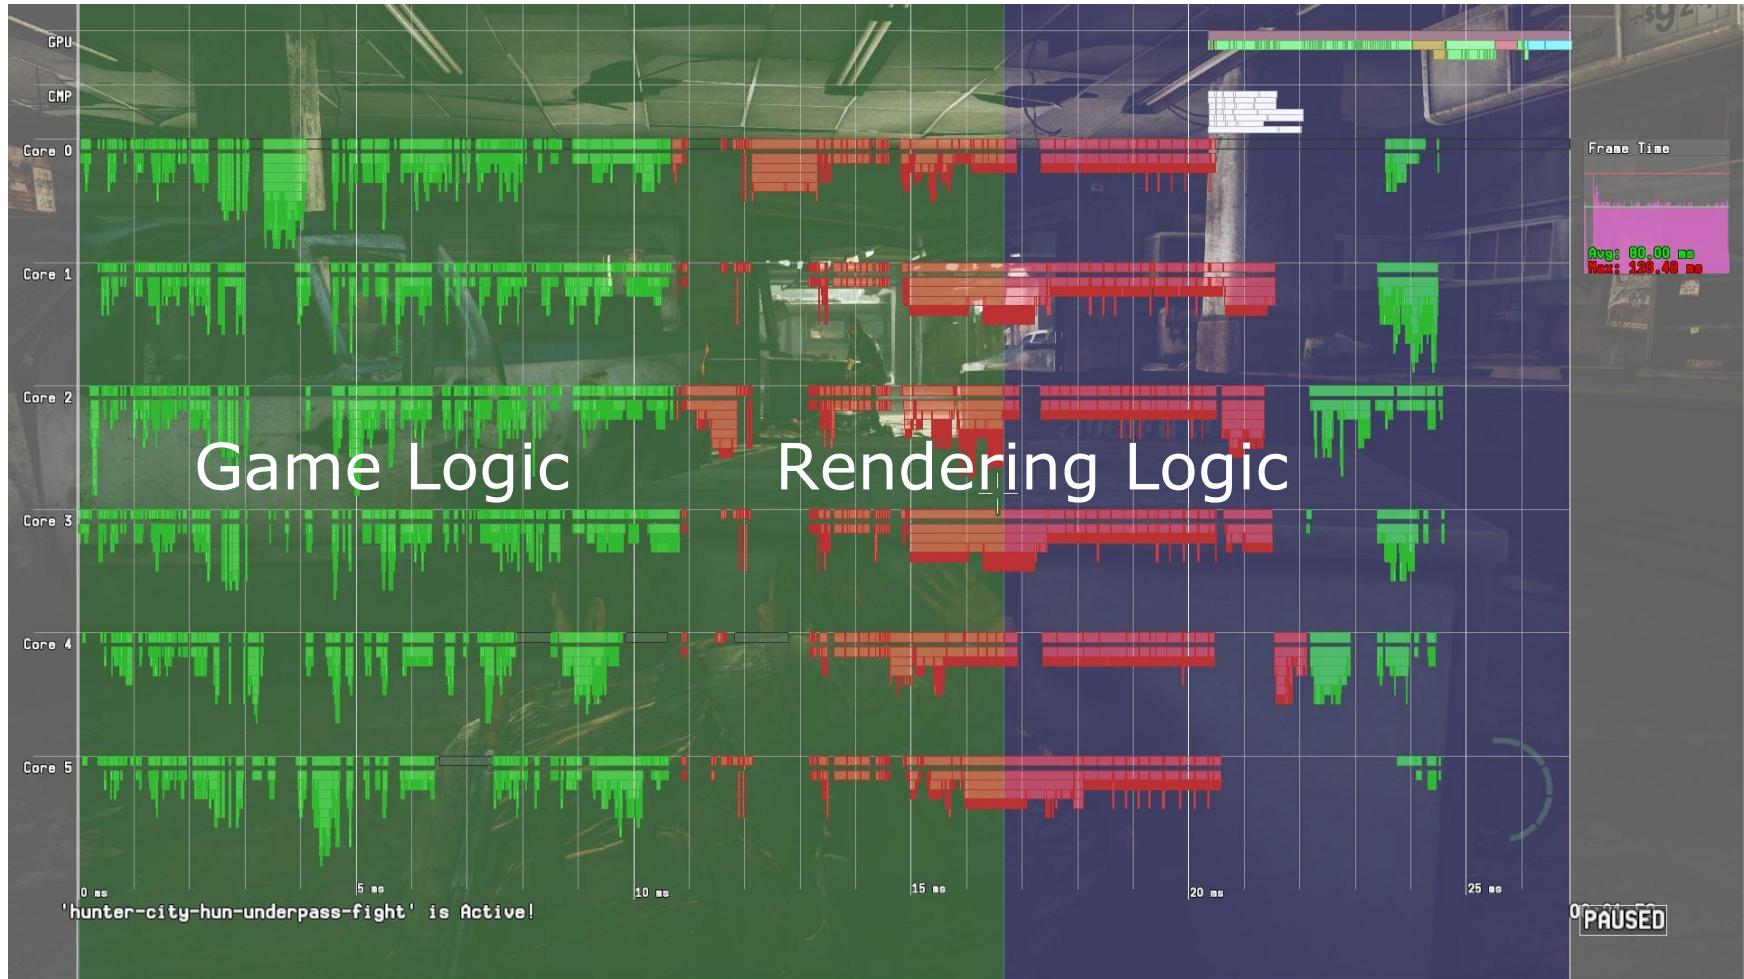


# Harsh realization

- Most systems are now jobified.
- GPU optimizations are going really well
- CPU bound
  - Critical path on CPU takes ~25 milliseconds.
- Lots of idle time on CPU cores
- At this point we are a little over 2 months away from shipping









# Is it possible to reach 60 fps?

- We had ~100 ms worth of work to do on the CPU for a single frame
- If we can manage to fill all the cores 100% of the time, then we can make it run at 60 fps.
  - $16.66 \text{ ms of work} * 6 \text{ cores} \rightarrow \sim 100 \text{ ms of work}$
- Ok, theoretically possible... but how do we do that?



# Let's cut it in half!!!



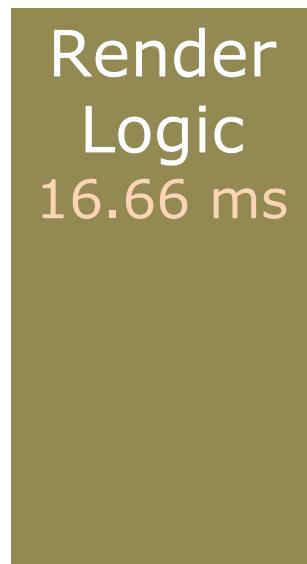
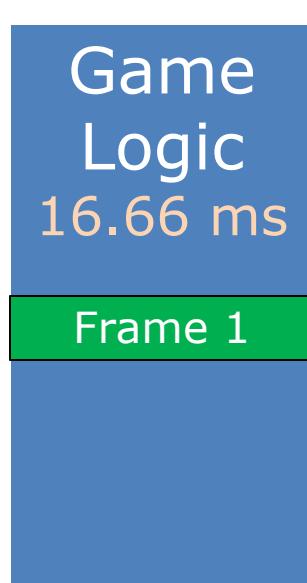


# How?

- Current design
  - Game logic runs first
  - Rendering logic and command buffer generation runs after
- New design
  - Run game and render logic at the same time
  - ***BUT processing different frames!***

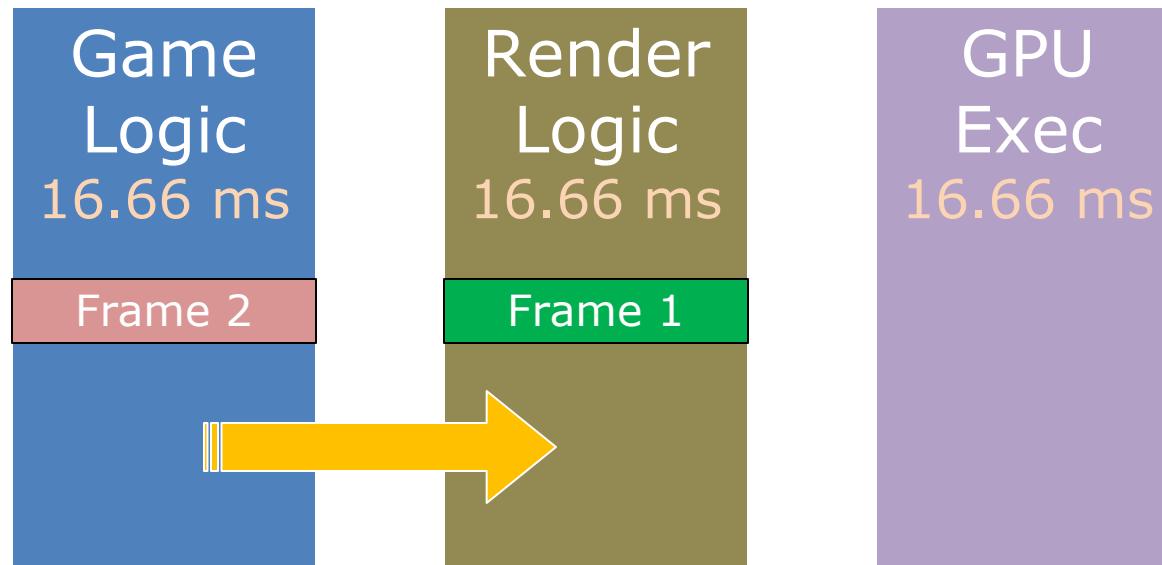


# Engine pipeline – Feed forward

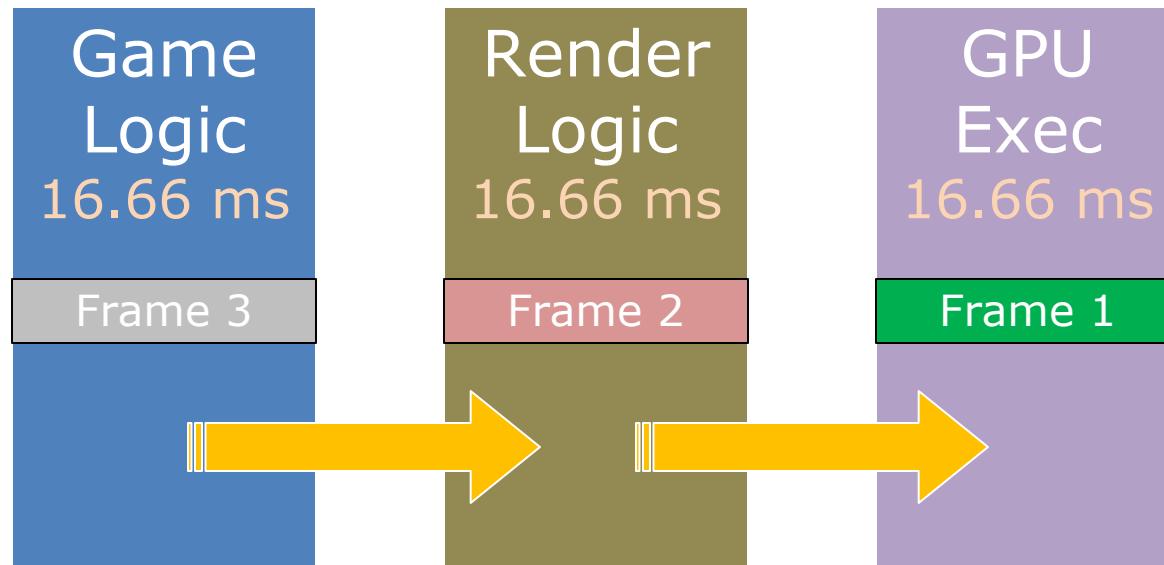




# Engine pipeline – Feed forward



# Engine pipeline – Feed forward



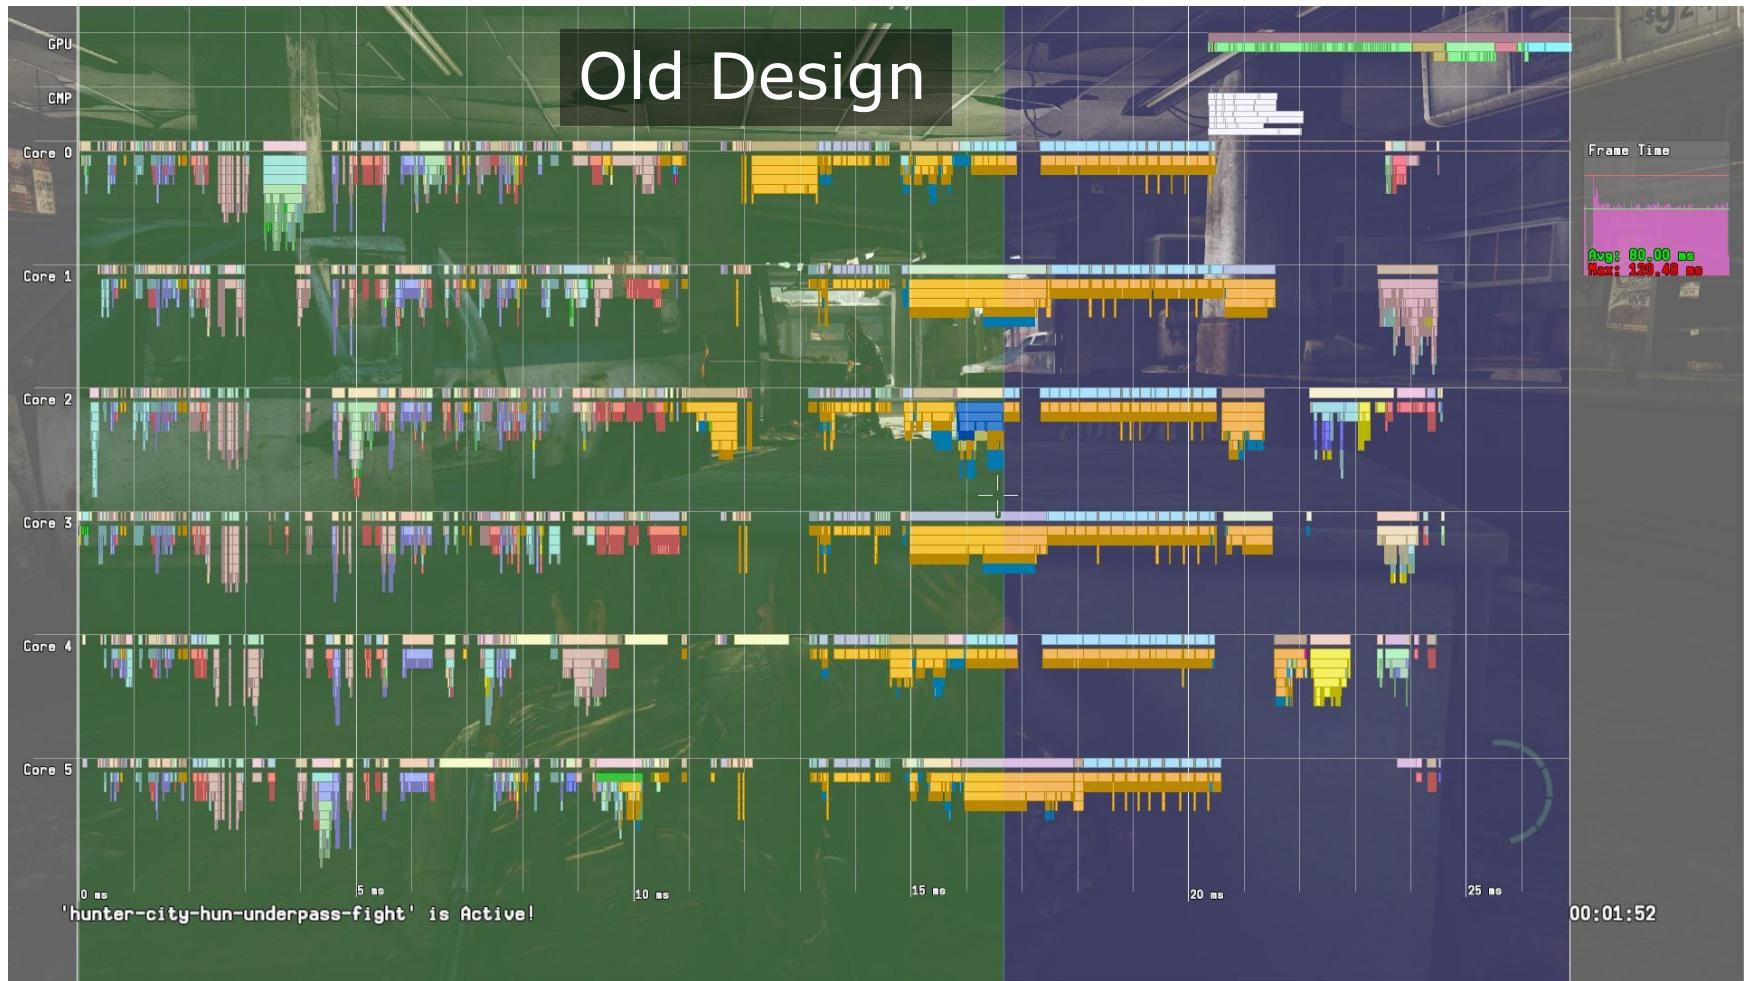


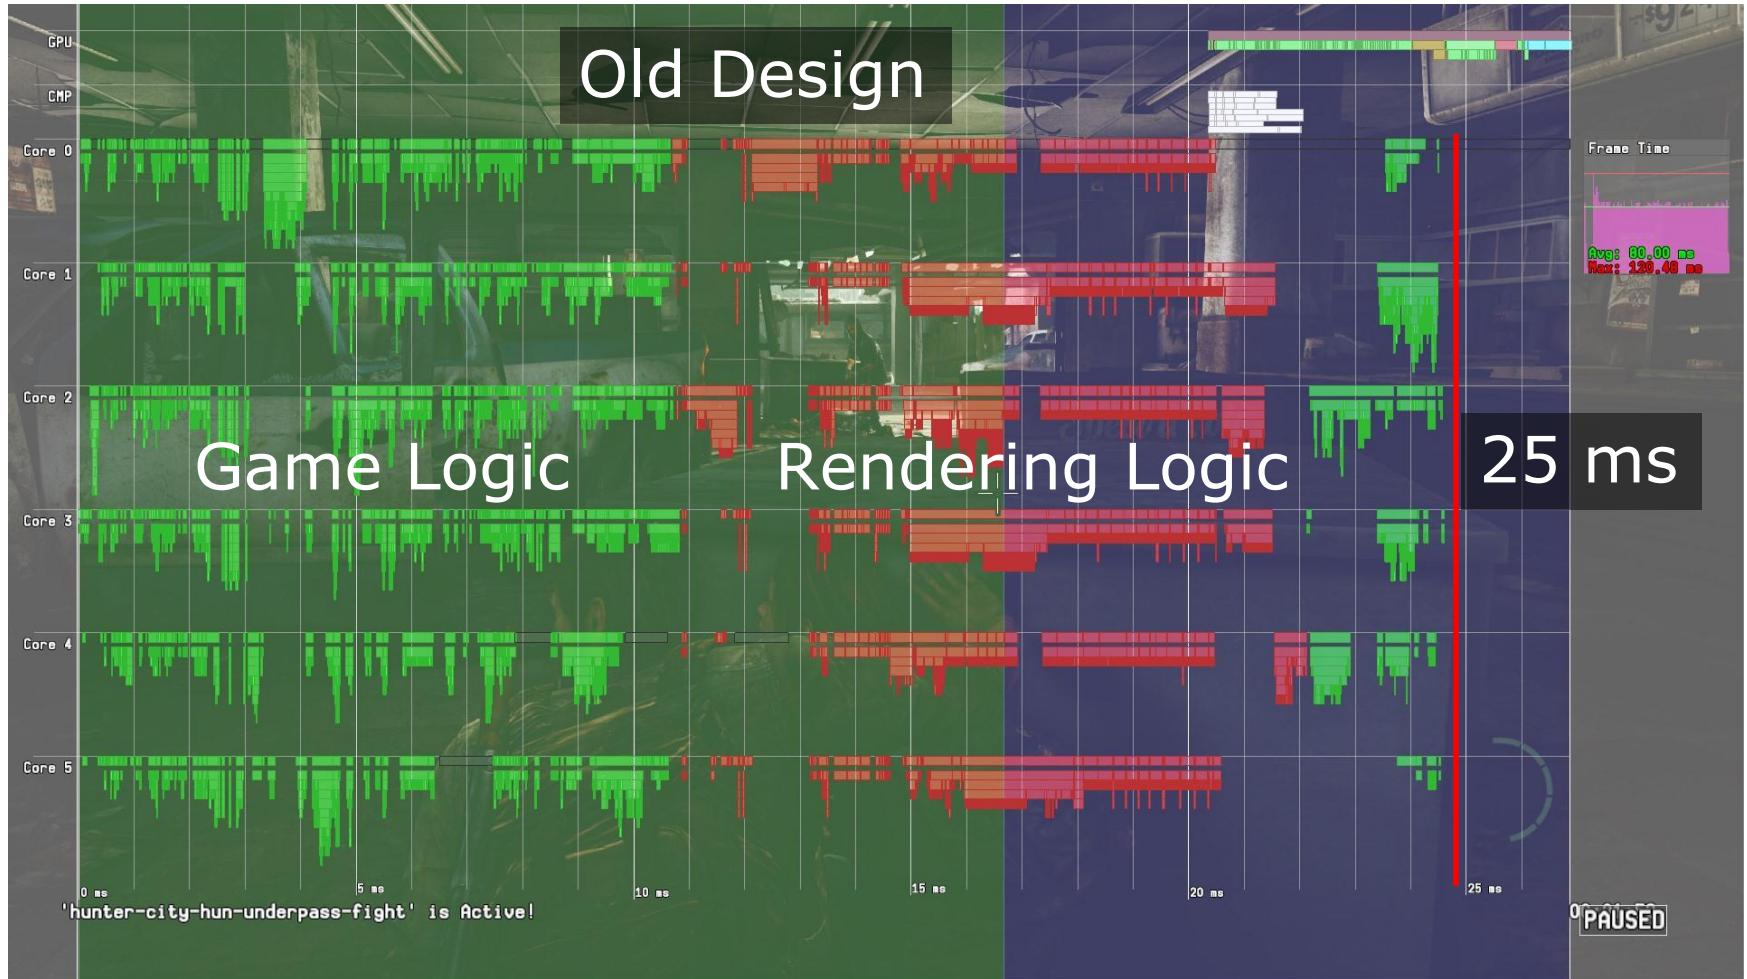
# Frame centric design

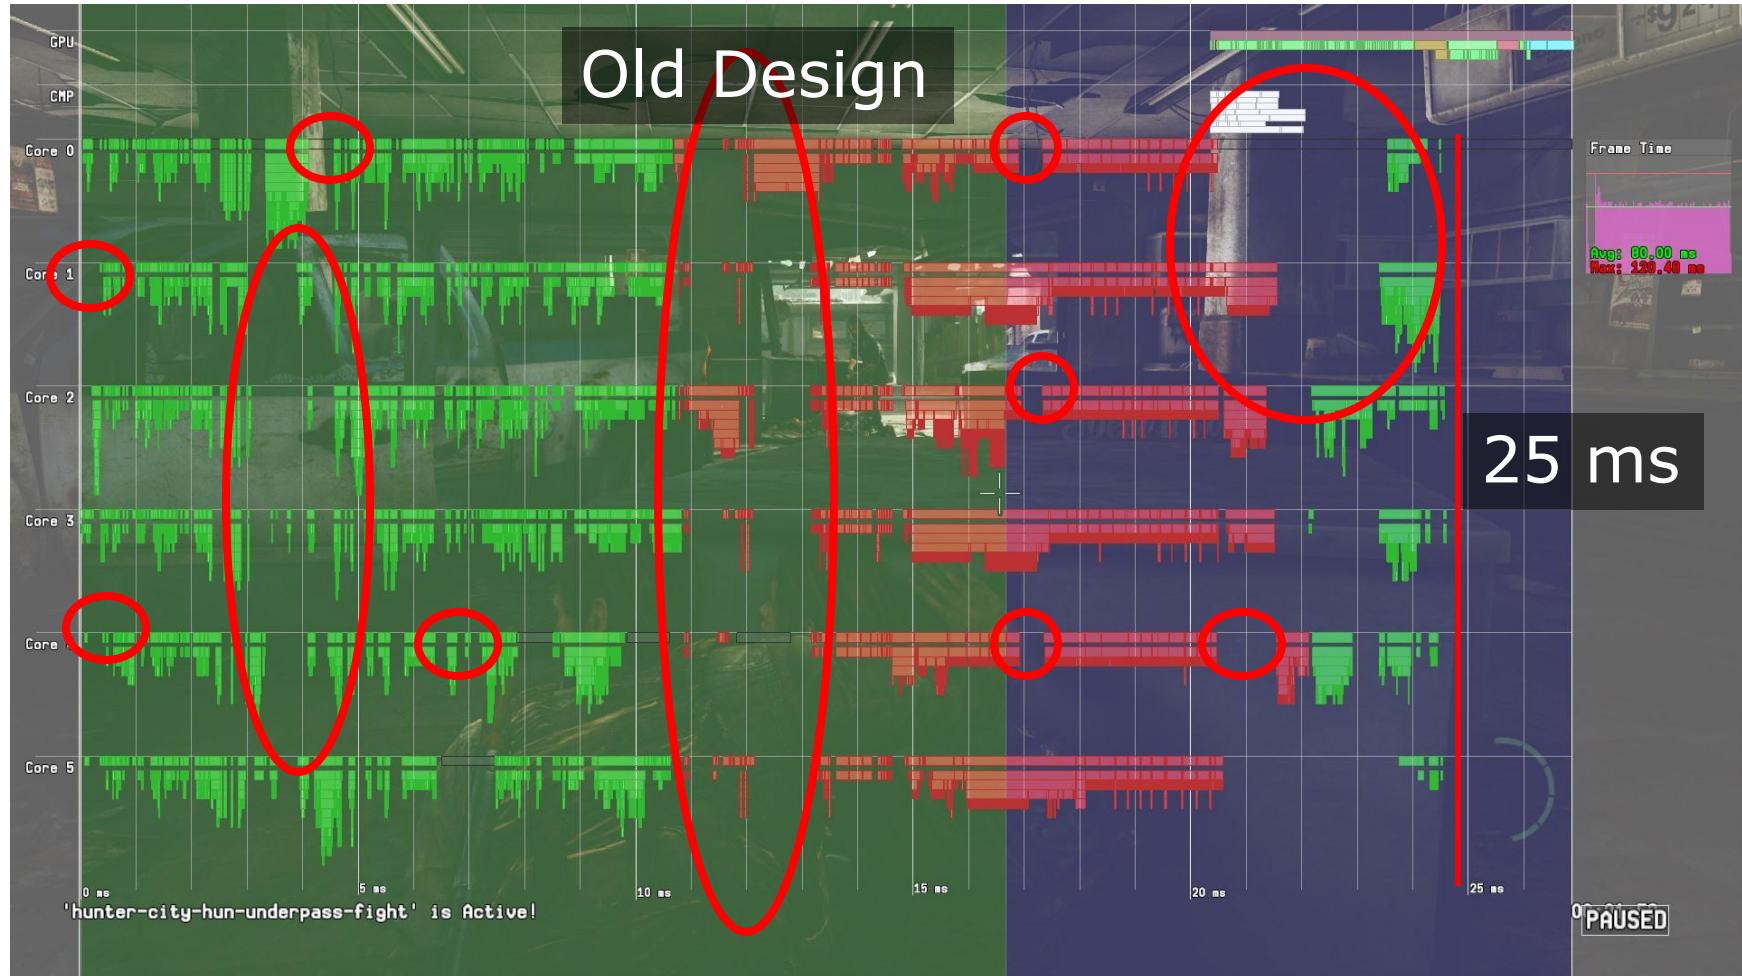
- Each stage is running completely independent
  - No synchronization required
- A stage can process the next frame right away
- Complexity in engine design is simplified
  - Very few locks needed due to parallelism
  - Locks are only used to synchronize **within** heavily jobified stage updates



# Old Design

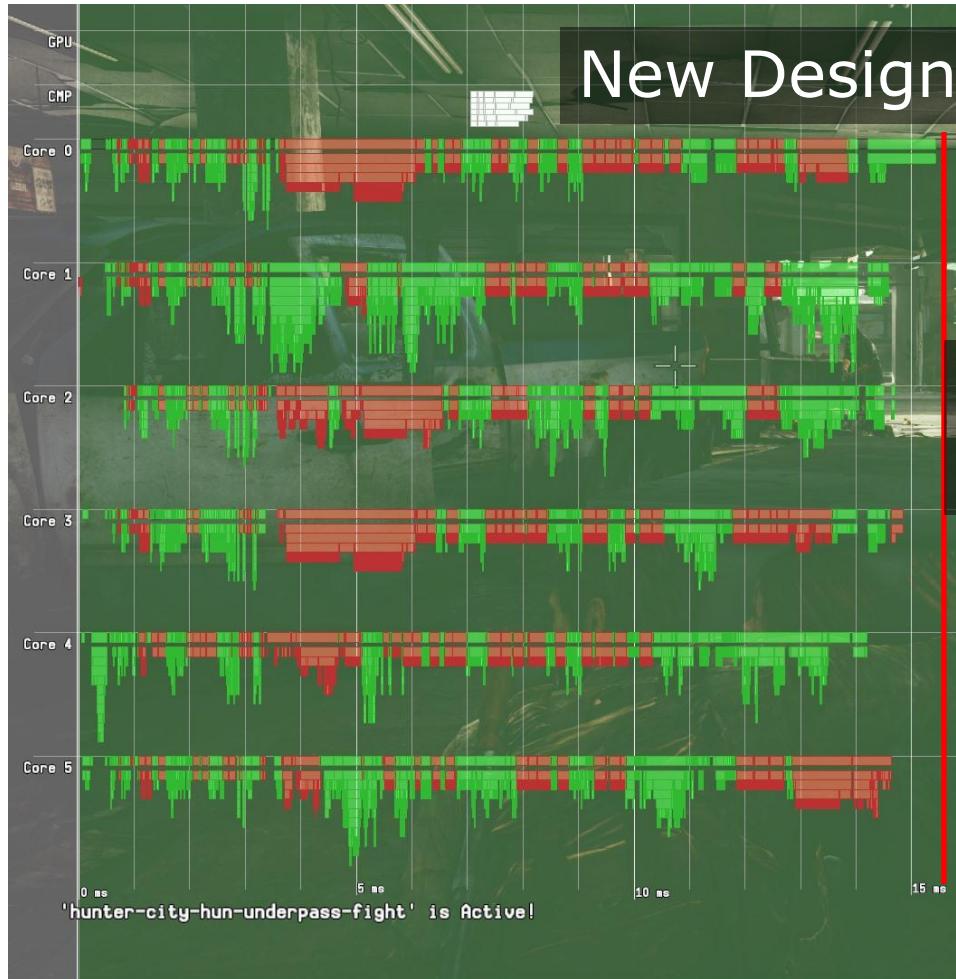








# New Design



15.5 ms!!!  
Ship it!





# Memory allocation on PS3

- Linear allocators
  - Contiguous block of memory + offset in block
  - Allocation is as simple as updating the offset
  - 'Free' all allocations by setting the offset to 0
- Single/Double/Triple frame
  - At the beginning of the frame we rotate to a new block and set the offset to 0
- Works great when you process ONE frame at a time



# Multi-frame difficulties

- “When should I free my memory?”
  - Memory is consumed in a later asynchronous stage
- “I need a new buffer every frame.”
  - “How many buffers do I need to rotate between to avoid memory stomping?”
- “Which delta time should I use?”
- Every programmer was solving the same problem
  - ... often incorrectly
- Hard to understand everything that is going on



# Frames, frames, frames...

- What is a 'frame' anyway?
  - Game logic frame, GPU frame, display frame?
- Memory lifetime
  - When do I safely free memory?
  - Can we double/triple buffer memory?
    - What does the even mean now?
- Need a new way of thinking about frames



# Our definition of a frame

- “A piece of data that is processed and ultimately displayed on screen”
- The main point here is ‘piece of data’.
- It is NOT a length of time.
- A frame is defined by the stages the data goes through to become a displayed image



# Introducing: FrameParams

- Data for each displayed frame
  - One instance for each new frame to eventually be displayed
  - Sent through all stages of the engine
- Contains per-frame state:
  - Frame number
  - Delta time
  - Skinning matrices
- Entry point for each stage to access required data



# FrameParams...

- Uncontented resource
  - No locks needed as each stage works on a unique instance
- State variables are copied into this structure every frame
  - Delta time
  - Camera position
  - Skinning matrices
  - List of meshes to render
- Stores start/end timestamps for each stage
  - Game, Render, GPU and Flip

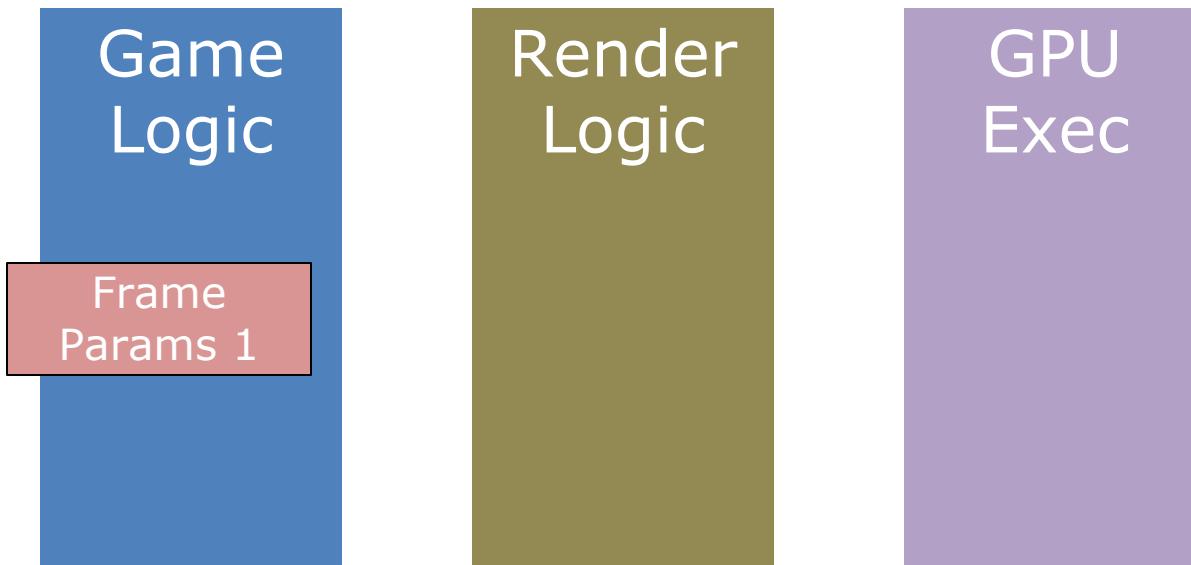


# FrameParams...

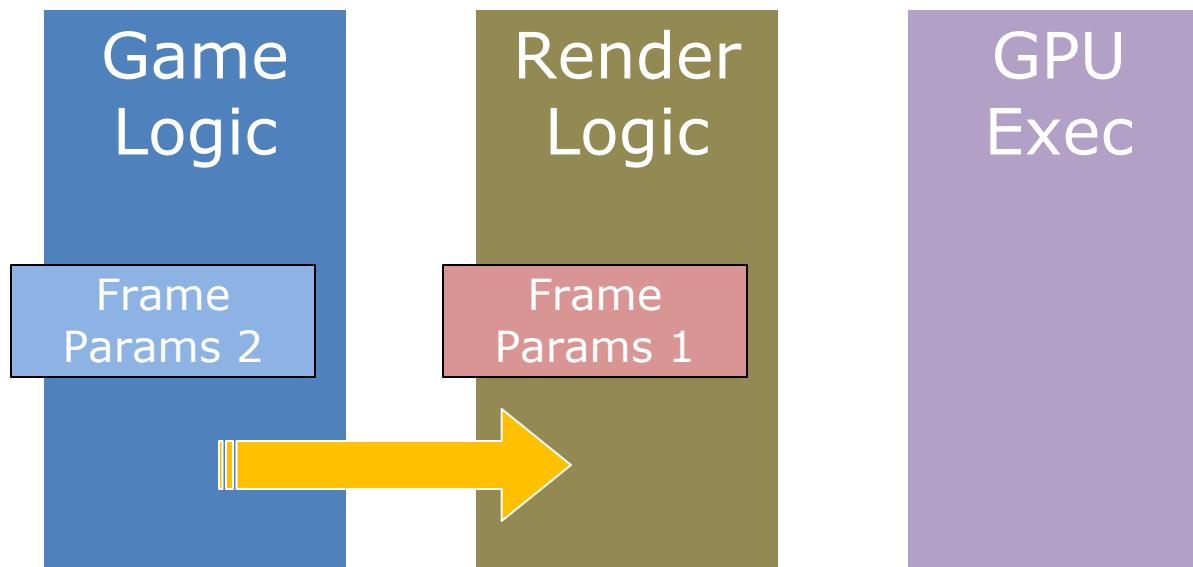
- Easy to test if a frame has completed a particular stage
  - HasFrameCompleted(frameNumber)
- Memory lifetime is now easily tracked
  - If you generate data to be consumed by the GPU in frame X, then you wait until HasFrameCompleted(X) is true.
- We have 16 FrameParams that we rotate between
  - You can only track the state of the last 15 frames



# Engine pipeline – FrameParams

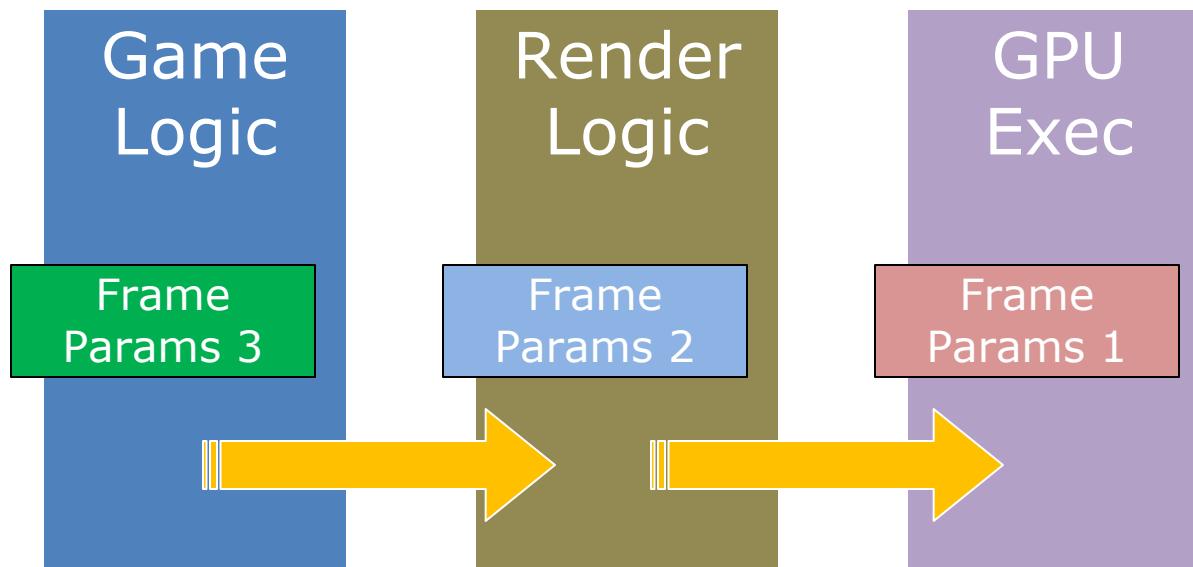


# Engine pipeline – FrameParams





# Engine pipeline – FrameParams





# Memory lifetimes

- Single Game Logic Stage (scratch memory)
- Double Game Logic Stage (low-priority ray casts)
- Game To Render Logic Stage (Object instance arrays)
- Game To GPU Stage (Skinning matrices)
- Render To GPU Stage (Command Buffers)
- ... for both Onion(CPU) and Garlic(GPU) memory!



# Running out of memory

- Many different linear allocators
- Many different life times
- All sized for worst case
- We never hit worst case for all allocators at the same time
- 100-200 MiB of wasted memory

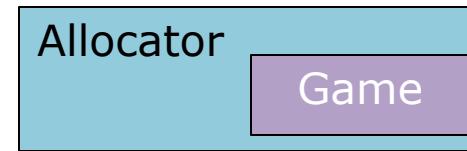
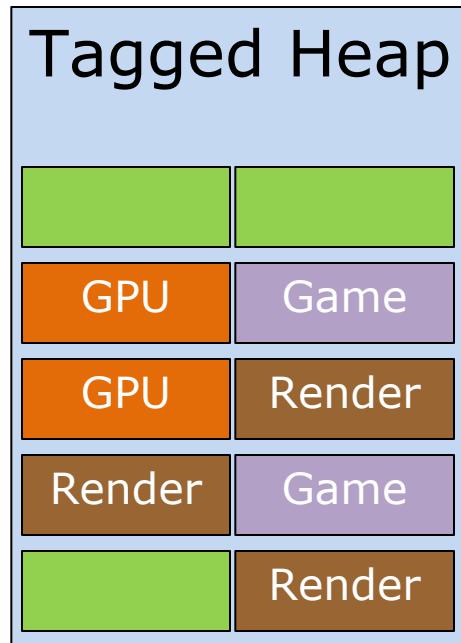


# Introducing: Tagged Heap

- Block based allocator
  - Block size of 2 MiB
  - 2 MiB is a 'Large Page' on PS4 -> 1 TLB entry
- Each block is owned by a tag
  - `uint64_t`
- No 'Free(ptr)' interface
- Free all blocks associated with a specific tag

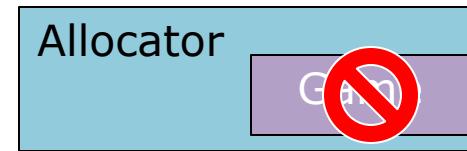
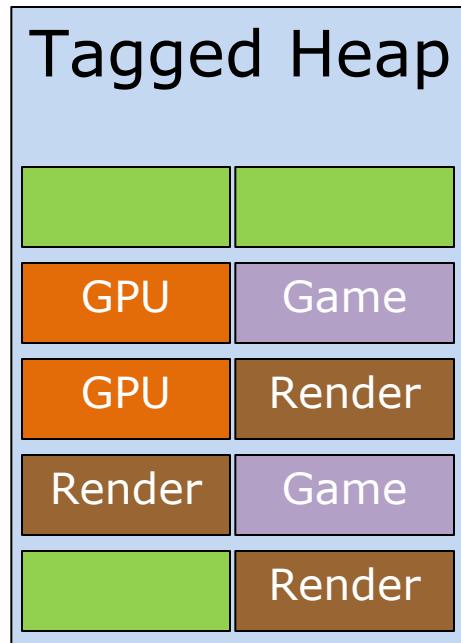


# Allocate from Tagged Heap



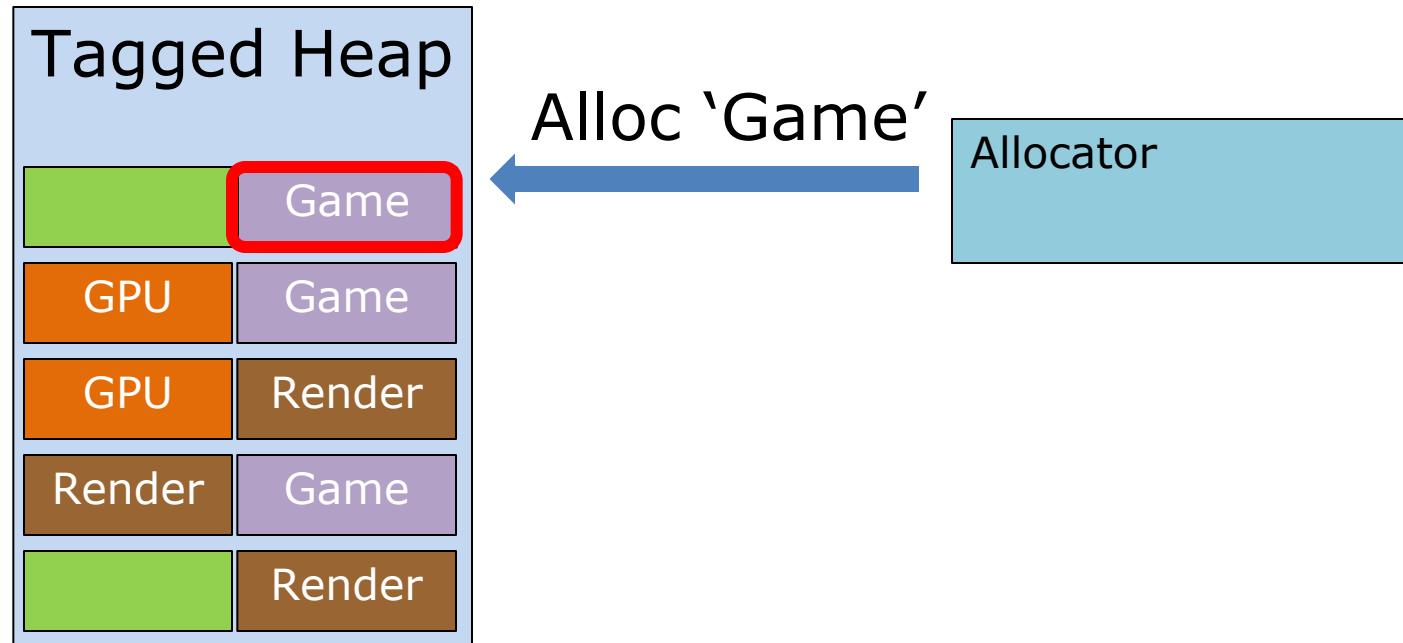


# Allocate from Tagged Heap



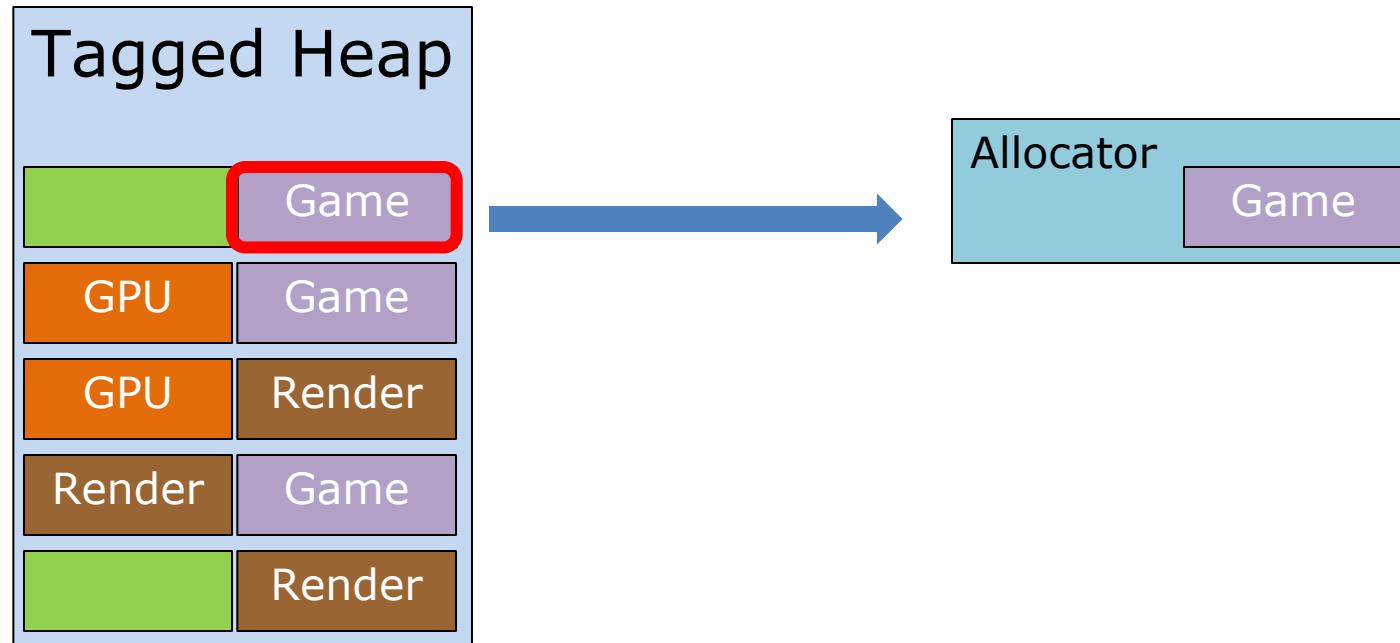


# Allocate from Tagged Heap



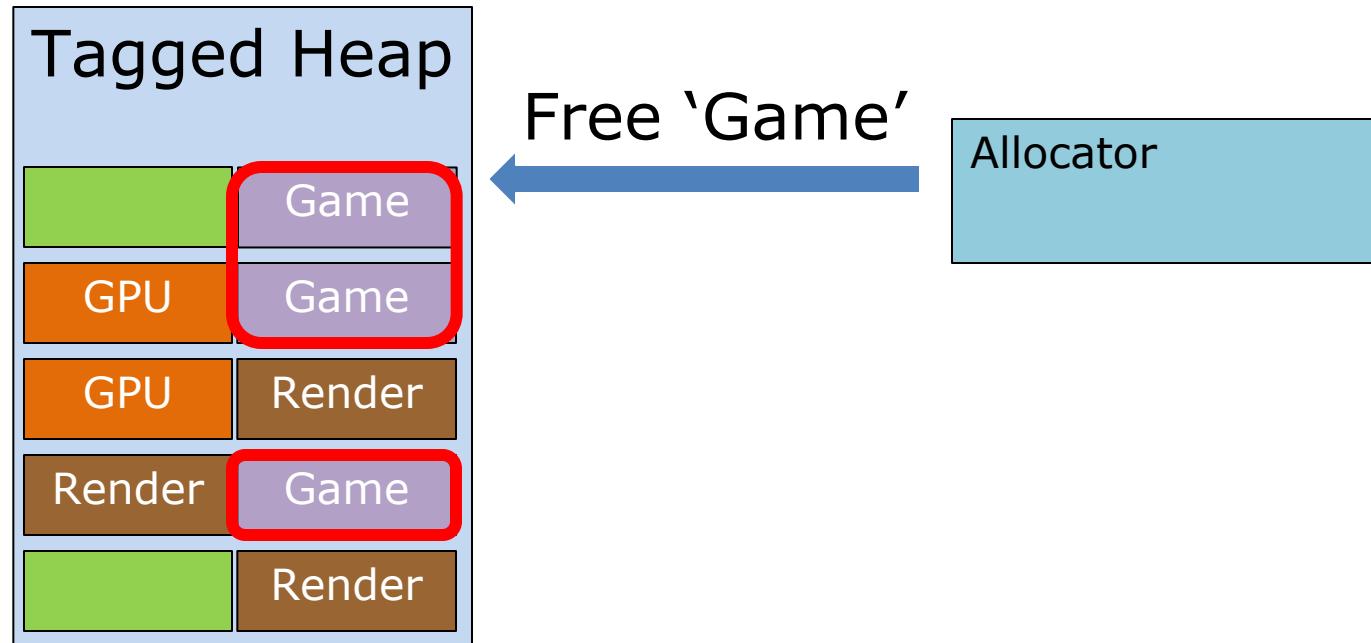


# Allocate from Tagged Heap



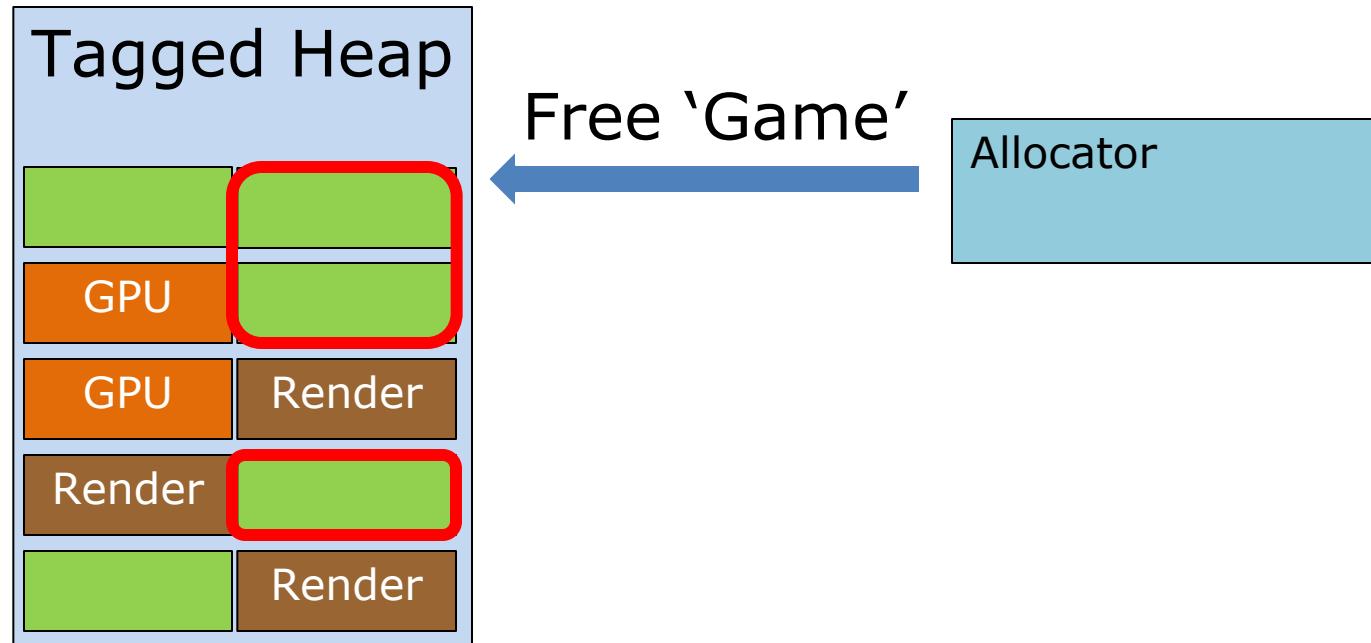


# Bulk free





# Bulk free



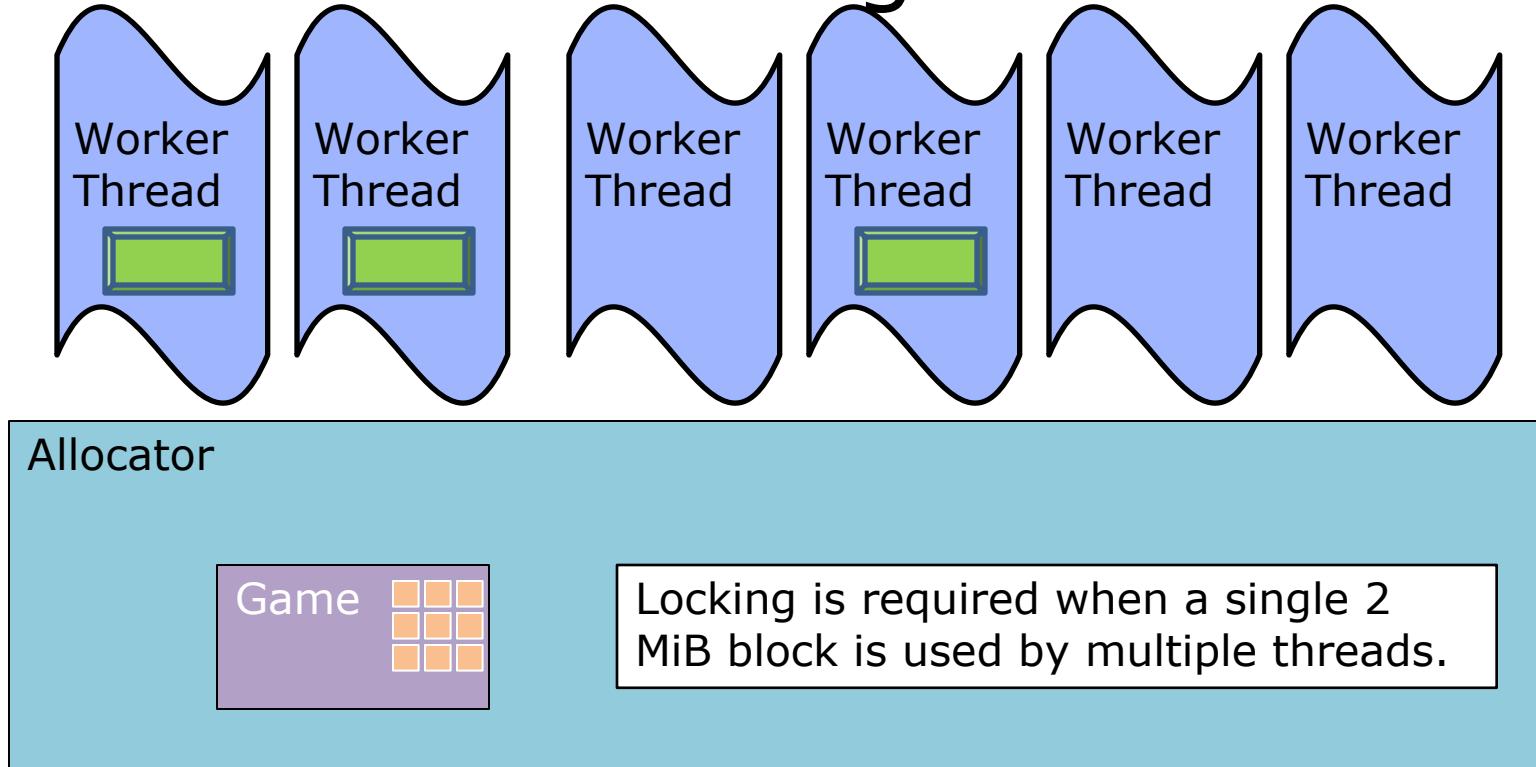


# All allocators use the tagged heap

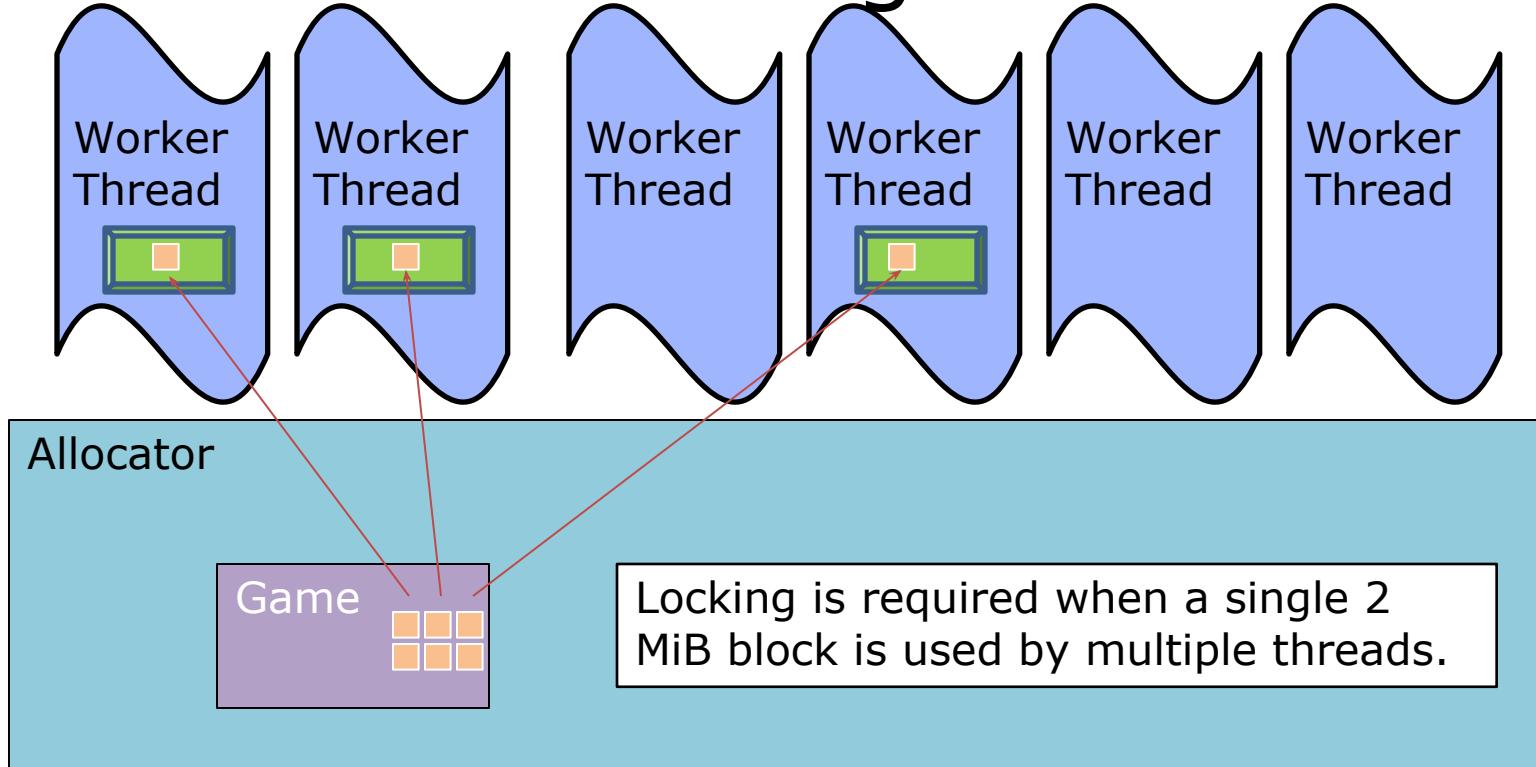
- Allocate 2 MiB block from shared tagged heap and store locally in the allocator
- 99% of all of our allocations are less than 2MiB
  - Larger than 2 MiB allocations get consecutive 2 MiB blocks allocated from the tagged heap
- Make allocations from this local block until empty
- Sharing a common pool of blocks like this allows for dynamic sizing of allocators



# Allocator with single block



# Allocator with single block



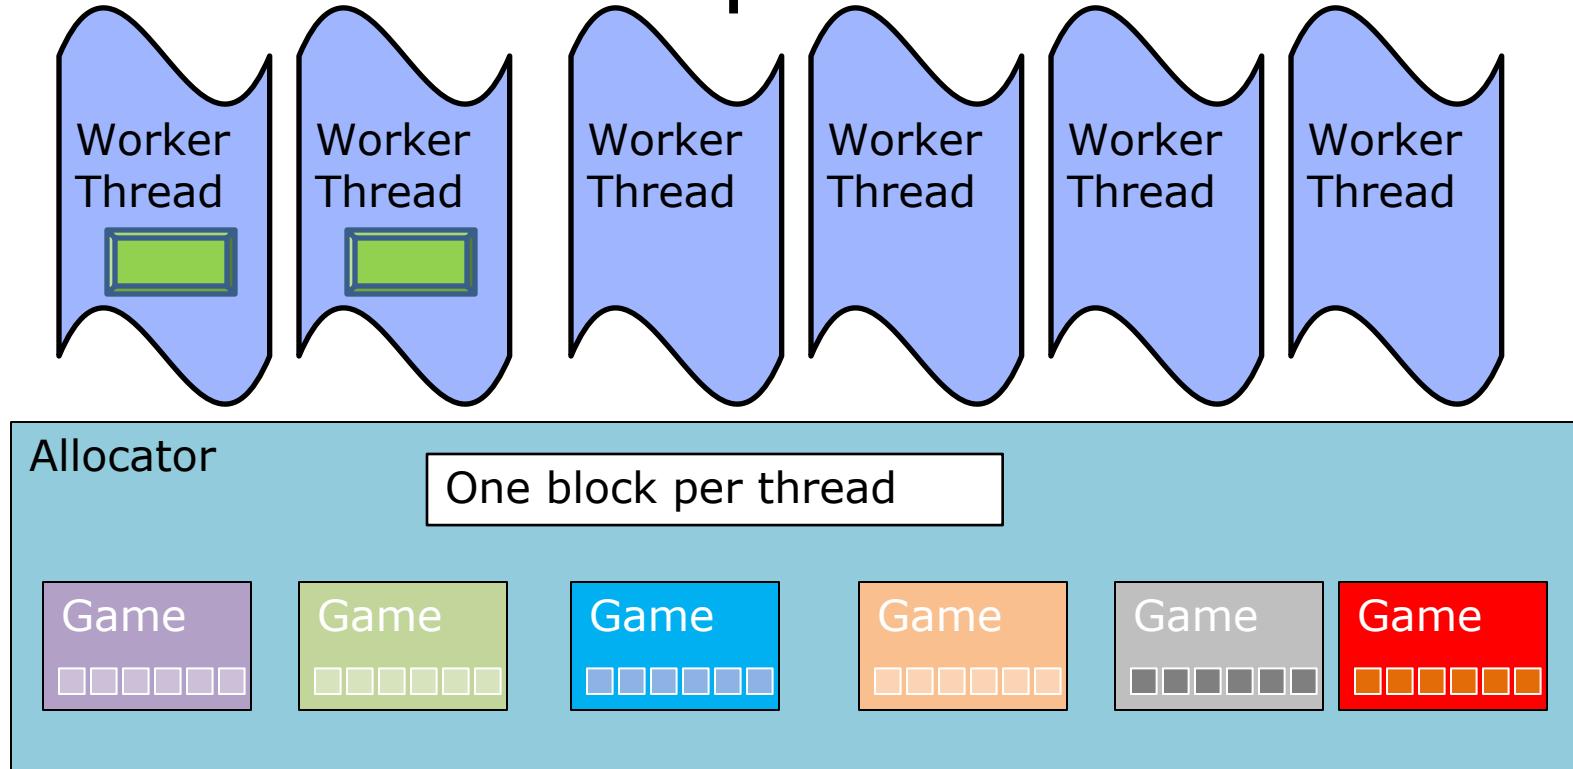


# Faster...

- Store one 2 MiB block per worker thread in the allocator
  - We use worker thread index to select which block to use
- All allocations on that thread go to that thread's block
  - No contention
- No locks required for 99.9% of memory allocations
- High volume, high-performance allocator

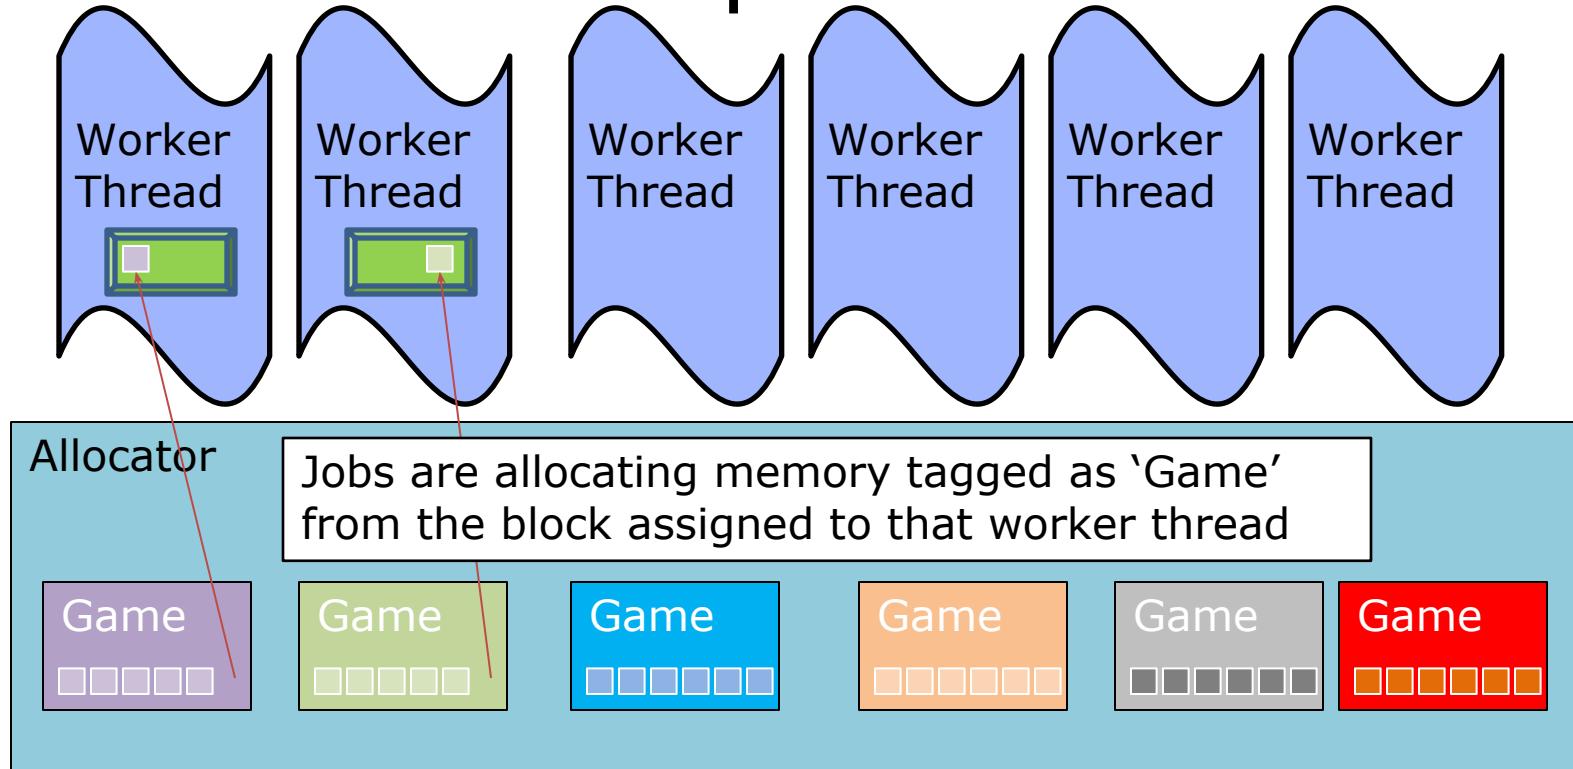


# Allocator with per-thread blocks



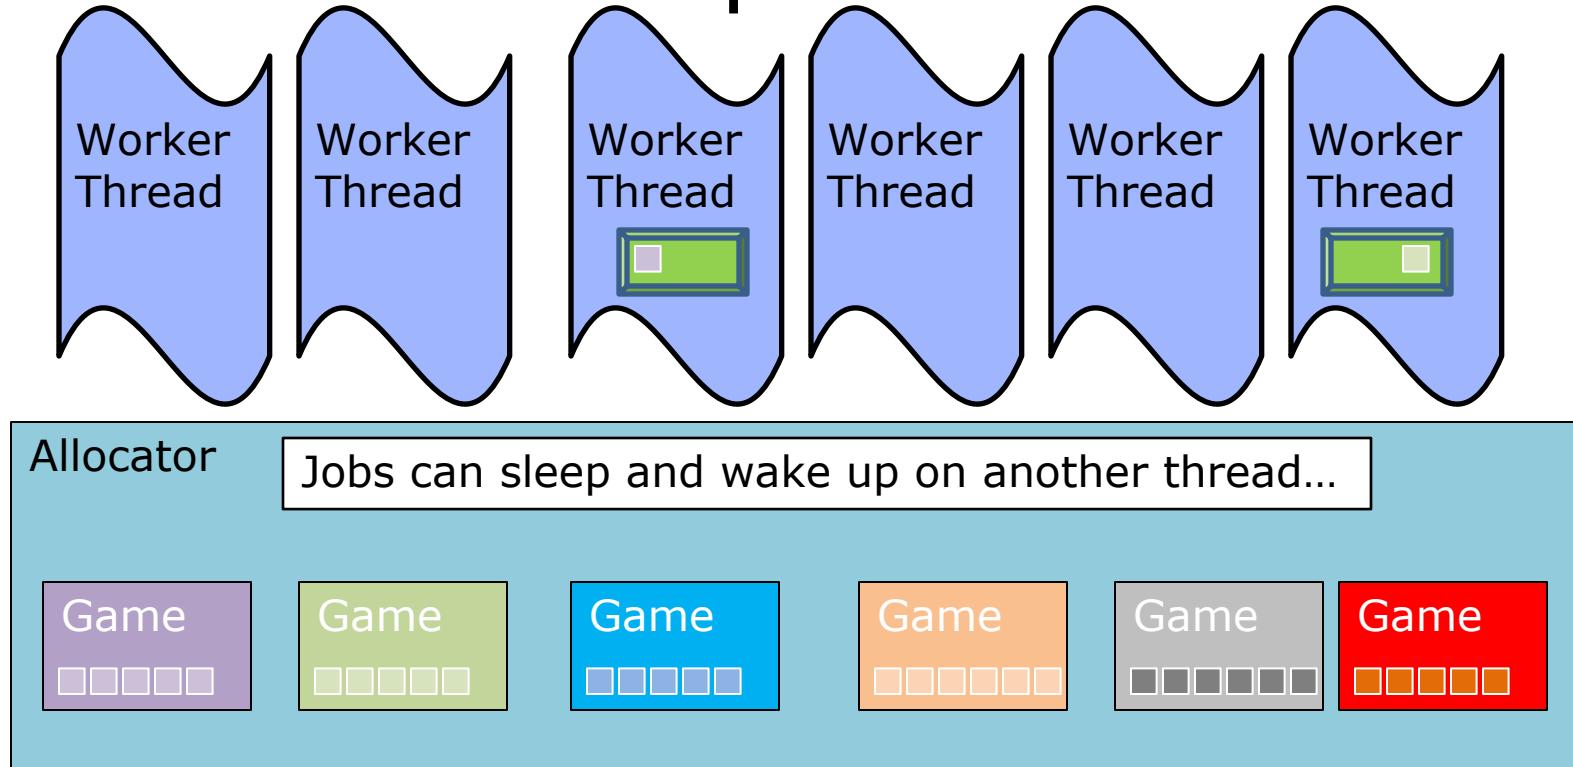


# Allocator with per-thread blocks



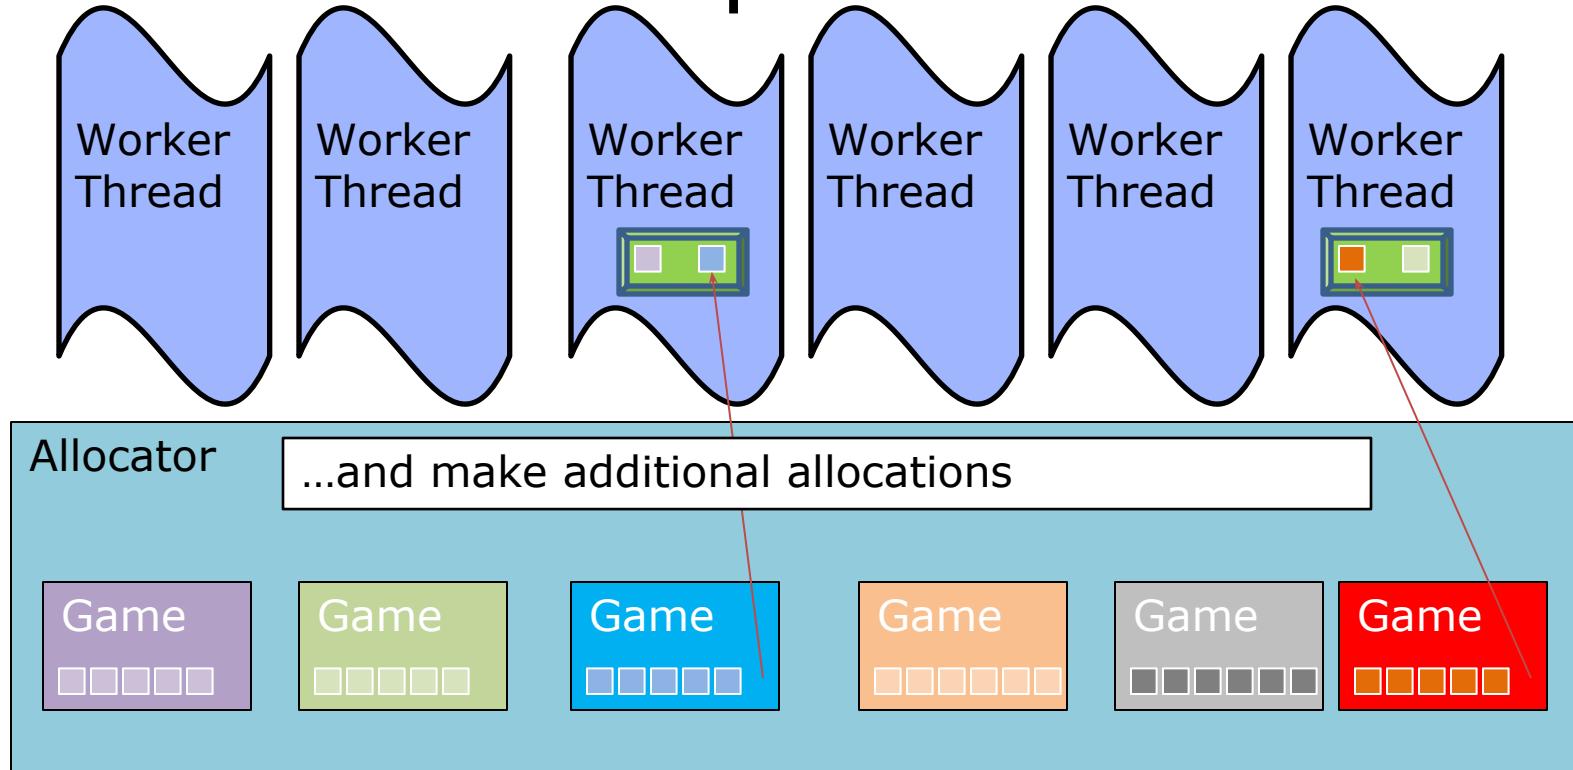


# Allocator with per-thread blocks



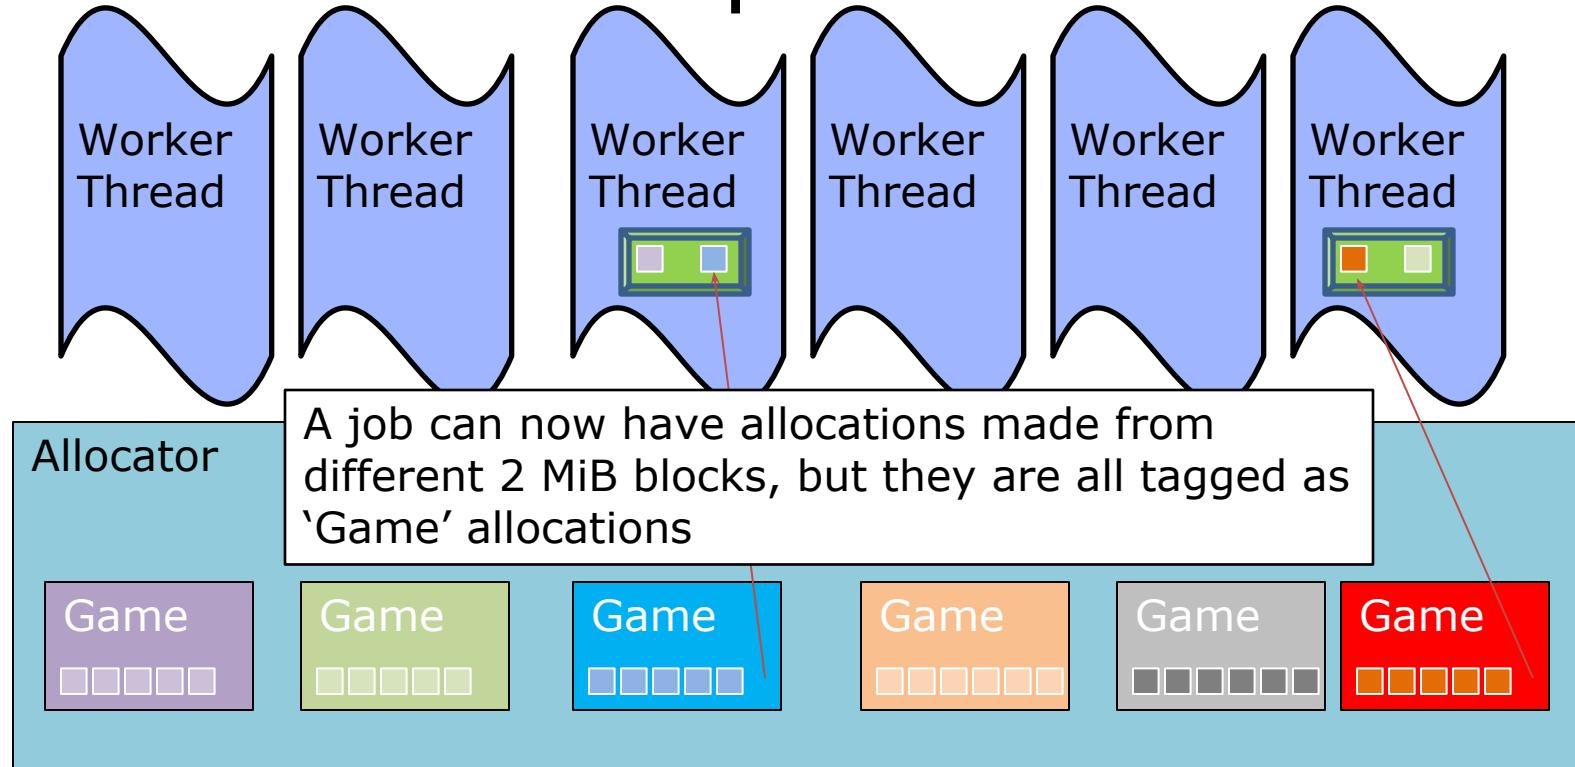


# Allocator with per-thread blocks





# Allocator with per-thread blocks





# Summary

- Fibers are awesome
- Frame-centric design simplifies your engine
  - Using something like FrameParams greatly simplifies data lifetime and memory management
- Tag-based block allocators are great when dealing with a multi-frame engine design



# Thank you!

# Questions?

Christian Gyrling

[christian\\_gyrling@naughtydog.com](mailto:christian_gyrling@naughtydog.com)