

The Technology of Jak & Daxter

Presented By Stephen White
Programming Director
Naughty Dog, Inc.

Naughty Dog is a Wholly Owned Subsidiary of Sony Computer Entertainment America, Inc.

Jak & Daxter: The Precursor Legacy

- Published by Sony Computer Entertainment for the PS2 in December 2001
- 35 full-time developers
- 3 years to develop
 - 1 year of initial development
 - 2 years of full production

Commercial Tools

- Operating systems used: Windows NT, Windows 2000, & Linux.
- Software used: Allegro Common Lisp, Visual C++, GNU C++, Maya, Photoshop, X Emacs, Visual SlickEdit, tcsh, Exceed, & CVS.

Internal Tool Development

- The importance of tools.
- Large amount of data that needs to be processed.
- Data formats and data streams need to be able to be changed fairly easily.
- Poor planning can result in a large loss of productivity.

The DB

- Boxed (tagged) data – could identify itself.
- The basic DB class was a type of generic container that could hold any number of other DB classes.
- Reading, writing, inspecting, copying, iterating, adding, and deleting can be done generically.
- New data can be added easily.
- Data is easily passed thru the pipeline.
- Any data file could be inspected by our generic db inspection tool.

Script Driven Tools

- Most of our tools were script driven.
- Generic script parser:
 - Could parse C syntax.
 - Had some additional macro support.
- Easy to add new features and options to the tools.
- Not the best interface for artists.
- Caused some undesired build dependencies.

From Maya To DB

- Only two stand-alone tools linked with Maya:
 - mbo (Maya Build Object)
 - mbl (Maya Build Level)
- Both tools did a fairly straightforward conversion from a Maya scene to our DB data representation.
- The other stand-alone tools used our DB data (not Maya's data representation).

From Maya To DB (cont)

- This insulated the other tools from changes in Maya, peculiarities of Maya, and having to link with Maya's libraries.
- Also allowed us to slightly cook and streamline the data to just what we were interested in.
- Kept data accessing nice and consistent for the other tools.

Additional Tools in the Pipeline

- buildactor – processed animations and foreground geometry, and outputted run-time data.
- stripdb – stripped background geometry.
- visdb – precomputed visibility to aid in run-time occlusion culling.
- buildlevel – final processing of level data, and outputted run-time data.

Additional Tools in the Pipeline (cont)

- tcomp – compiled our texture pages.
- texinfo – Windows tool used by artists to manage textures and texture pages.
- Various Maya plug-ins and MEL scripts to aid in modeling, lighting, texturing, instancing, creating/editing tfaces, placing shrubs, etc..
- Various other tools – Viewing db files, examining databases, processing sounds, debugging, etc..

Tool Things We Did Wrong

- Didn't use a PC farm to process levels and actor animations.
- Overburdened our network.
- Poorly constructed data dependencies.
 - Small changes could cause needless and lengthy reprocessing of data.
 - Relied too much on file time/date stamp checking, which turned out to be extremely slow.

Tool Things We Did Wrong (cont)

- Underestimated the difficulty of managing the audio, and had insufficient audio tool support.
- Processing FMA sequences was difficult, time-consuming, and error prone.

Background vs. Foreground

- We typically refer to things that are mainly static as “background” (i.e., ground, cliffs, rocks, trees).
- We typically refer to things that can move around as “foreground” (i.e., an enemy).
- This is a seemingly arbitrary distinction, but affected how our data and renderers were constructed.

Background Elements

The background features a dark blue gradient with a large, light blue curved shape on the right side, creating a dynamic, abstract design.



Tface (Tessellating Face)

- Dynamically tessellates based on the distance to the camera.
- High detail up close, low detail far away.
- Composed of two primitive types:
 - tface quads
 - tface tris


Tface Quads

- Essentially a grid of polygonal quads.
- 1, 2, or 4 quads wide by 1, 2, or 4 quads tall (i.e., 1x4, 2x2, 4x2, etc).
- Modeled as a collection of quads in Maya.
- Custom Maya plugins allowed tface quads to be “constructed” from polygonal quads.
- Custom Maya plugins aided in setting/editing uv and corner color values.

Tface Quad Examples



Tface Quad Tessellation




4x4 (16 quads) reduces to a single quad.

93.75% reduction in detail!


Tface Tris

- 1, 2, or 4 subdivisions for each edge (i.e., 1x4x2, 2x2x4, 4x1x4, etc).
- Modeled initially in Maya as a single triangle, then tessellated by a plug-in into the required configuration.
- Custom Maya plugins aided in setting/editing uv and corner color values.

Some Tface Tri Examples



Tface Tri Tessellation



4x4x4 (16 tris) reduces to a single tri.

93.75% reduction in detail!

How Lawyers Become Rich

- Sony is in the process of patenting how exactly we did the tessellation of our tfaces.

Tfrag (Tessellating Fragment)

- Composed of a collection of tface tris & tface quads.
- Static, non-instanced (unique) background.
- The tfaces of a level were broken up by the tools into tfrags based on stripping and Vector Unit memory size constraints.
- Tfrag geometry indexed a 1024 color palette for corner color lighting.

Tfrag Renderer(s)

- Used multiple renderers based on how a tfrag was to be tessellated (distance to camera).
 - Optimized renderers for various cases.
- Supported effects:
 - Opaque
 - Translucent
- Translucent tfrag was not sorted, so draw order could be wrong.
 - In practice, this wasn't as bad as you'd think.

Tfrag Collision

- Modeled as tfaces for convenience, but converted by the tools to tris.
- The artists could use either the geometry to render as the collision geometry or could build custom collision geometry as appropriate.
- Only one collision model was modeled for all tfrag, but the tools broke up the data into smaller pieces written out as a spatially subdivided tree.
 - Each node of the tree had a bounds sphere and up to 8 child nodes.

TIE (Tfrag Instance Engine)

- Instanced: one model (the prototype) drawn multiple times, using a different matrix and lighting per instance.
- Prototype modeled as tface tris & quads.
- Placed in a Maya scene file using ref nodes.
 - Ref nodes could be viewed expanded or collapsed using custom Maya plug-ins.
- Each instance was lit individually, but was then mapped to a single 128 color palette per prototype.

TIE Renderer(s)

- Essentially instanced tfrag, but uses specialized renderers.
- Supported effects:
 - Opaque
 - Translucent
 - Environment mapped (thru the Generic Renderer)

TIE Collision

- Modeled identically to tfrag collision, but each TIE prototype had its own collision mesh (or no collision at all).

Shrubs

- Used for more than just plants.
- Fast renderer for rendering small details.
 - Data format closely matched what GS used to draw.
- Instanced (one model drawn multiple times).
- Prototype modeled as tris and quads.
- Placed in a Maya scene file using ref nodes.
 - A special Maya plug-in could also be used to “randomly” populate an area of ground with shrubs.
- Each instance had a color multiplier used for lighting.

Shrubs (cont)

- Supported effects:
 - Opaque
 - Translucent
- Could become a single quad “poster” in the distance (far level of detail).
- Becomes translucent and fades away in the distance.
- An instance could sway in the breeze using a manipulation of its transformation matrix.

Background Lighting

- The background was pre-lit up to 8 times using Maya.
 - An artist would light the level, then save the lighting information into a level color database.
 - The 8 sets of colors were typically used to achieve our “time of day” lighting.
 - At run-time the sets of colors could be cross-faded to create the appropriate lighting for the current time of day.

Background Lighting (cont)

- Tfrag vertices indexed a 1024 color palette that was generated using 8 tfrag time of day palettes.
- TIE vertices indexed a 128 color palette that was generated using the prototype's 8 time of day palettes.
- Other effects, such as flickering lights, were also done using the 8 palettes.

Misc Background Renderers

- Sky renderer.
 - Looking at the inside of a pyramid.
- Ocean renderer.
 - A giant tessellating grid.
 - Nearby polys were translucent, while farther polys were opaque.
 - Environment mapped.

Actor Placement

- Placed using specially named locator nodes.
- Assigned a unique name (i.e., grunt-67) based on Maya DAG path and scene file.
- Assigned attributes in a hand-typed level config script.

Curves

- Linear and second degree curves.
- Parented under an actor node, and saved as part of that actor's instance data.

Nav Meshes

- Used for navigating enemies around our complex 3D world.
- Modeled as a relatively 2D mesh composed of tris.
- Described in detail in Game Programming Gems 3, published by Charles River Media, Edited by Dante Treglia.

Volumes

- Modeled as polygonal meshes in Maya.
- Used for a variety of things, including visibility, water detection, etc.

Foreground Elements

The background is a solid dark blue. A thin, light blue curved line starts from the top left and arcs towards the right. On the right side, there is a light blue triangular shape pointing towards the center, partially overlapping the dark blue background.

Merc

- Used for enemies, platforms, FMA, etc.
- Modeled as polygonal meshes composed as tris or quads, and bound (weighted) to joints.
- Mesh deformed by joints using a max of 3 joint influences per vertex.
- Optionally vertices could also be deformed using blend shapes.

Merc (cont)

- Used separate models for different levels of detail. Auto switched between models based on the distance to the camera.
- Becomes translucent in the distance and fades away.
- Also used a special eye renderer to composite animating lids, pupils, and irises.

Merc Blend Shapes

- Used mainly to do the highly animated facial expressions of the FMA characters.
- Multiple “blend targets” were created by the artists in order to generate various facial expressions.
- The animators could blend in the effects of multiple blend targets as desired.
- The artists used Maya’s blend shape support to both model and animate the blend shapes.

Merc Renderer(s)

- Supported effects:
 - Opaque
 - Limited, unsorted translucency
 - Environment mapping (thru the Generic Renderer)
 - UV scrolling
 - Ripple (multiple sin waves thru a grid to make waves of water)
 - Vertex deformation using joints.
 - Vertex deformation using blend shapes.

Foreground Lighting

- 3 dynamic lights used per merc object.
- Global “rgba multiplier” and “rgba add” used per object.
- Optionally, foreground could use time of day directional lighting information.

Foreground Collision

- Multiple collision meshes allowed per actor model.
- Each collision mesh was modeled as a polygonal mesh using tris or quads.
- Each collision mesh was transformed at run-time using a single joint (no vertex deformation).
- Optional collision spheres bound to individual joints we're also used, and were configured in code.

Shadows

- A simplified model was used to cast shadows.
- Mesh deformed by joints using a max of 2 joints influences per vertex.

Shadow Renderer

- A shadow volume was dynamically constructed at run-time based on extruding the transformed mesh, typically away from a light source.
- Once the extruded run-time volume was created, the shadow renderer using the classic shadow volume approach to create an alpha mask that was then used to darken the underlying pixels on the screen.

Misc Foreground Renderers

- Sprite (particle) renderer:
 - 2D sprites
 - 3D sprites
 - Distortion sprites (simplistic refraction effect)
 - Heat effects
 - Blurs.

The Generic Renderer

- Due to our large number of highly specialized renderers, certain effects were impractical to implement in each renderer.
- Other renderers (either foreground or background) could output to the Generic Renderer's format in order to do specialized effects or to handle computational or memory-wise expensive cases.
- Some examples:
 - Environment mapping
 - Scissoring (triangle subdivision)

The Generic Renderer (cont)

- The generic renderer was slower than the specialized renderers, but was much easier to extend to do various effects.

Enough About Foreground and Background

Visibility (Occlusion Culling)

- Visibility information about which fragments and actors could be seen from a given location was precomputed by the tools.
- A BSP tree was used to look up the appropriate visibility information at run-time based on the position of the camera.
- The visibility information was essentially a string of bits specifying “draw” or “don’t draw” for the various elements in the scene (I.e., tfrags, TIEs, shrubs, actor, etc).

Visibility (cont)

- The bit strings were stored compressed, and were decompressed as needed at run-time.

Defining Where an Actor was Visible

- Actor visibility was specified using a bounding box around everywhere that the actor could reasonably travel.
 - The actor bounding boxes could be auto-generated at run-time using both the actor nav meshes and by playing the game for awhile.
- An actor's visibility could be set specifically in the level config script for given locations.

Spooling Level Data

- Used two 10M data level buffers.
- A single game level was composed of one or more data “levels”.
- We used careful planning and layout to insure that no more than two data “levels” were visible at any one time:
 - The level the player was currently in.
 - The level the player was heading towards.
- As the player approached a new level, it was loaded into the other level buffer.

Spooling Level Data (cont)

- If the player changed his mind and headed towards a different level, then the level that was being loaded was abandoned, and a new level load was started.
- Used “load boundaries” to determine when to start loading or when to display a new data level.
- Used low resolution models to represent portions of other levels seen in the distance.

The Spooled Level Data

- The level data contained:
 - Background (geometry, collision, textures, visibility, etc).
 - Foreground (geometry, collision, textures, animations, etc).
 - Sound.
 - Level specific code.

Other Spooled Data (Not Part of the Level Data)

- Certain lengthy individual animations (i.e., Jak's bored/idle animations).
- FMA sequences: animation for multiple characters, and audio.
- Spooled audio for lengthy sound effects.

G.O.A.L.

(Game Object Assembly Lisp)

What is GOAL?

- GOAL is our custom compiler based on Lisp (well, actually Scheme).
- Practically all of the run-time code (approximately half a million lines of code) was written in GOAL.
- Only the IOP code and a small amount of kernel code was written in C.

GOAL Features

- Object-oriented language.
- Extremely simple syntax.
- Powerful macro capability, far superior to C's preprocessor or C++'s templates.
- Listener
 - Code could be executed live at a “listener”.
 - Code could be compiled, downloaded, and linked without interrupting gameplay.
 - Data structures could be inspected or modified live.
 - Rapid tuning and debugging.

Goal Features (cont)

- Non-preemptive, cooperative multi-tasking.
- The “suspend” operation.
- Stack variables were preserved across a suspend.
- Code could be written to flow naturally, rather than the typical C++ AI behavior schemes which use state based “update” callbacks that have to deduce what to do next.
- Extensive support for AI scripting.

Goal Features (cont 2)

- Unified set of assembly op-codes consistent across all five processors of the PS2.
- Seamless intermixing of high-level code with low-level (assembler) instructions.
- Register coloring when using assembler instructions.
- Code could be loaded as part of the level data.

Difficulties with GOAL

- Compiler took over a year to write, and was fully developed and supported by one person.
- We suffered with a buggy, and quirky compiler as GOAL was being developed.
- An unfamiliar language for programmers to use.
- Difficult learning curve.

Difficulties with GOAL (cont)

- Isolated from other developers. Couldn't use other people's tools or libraries.
- The compiler would periodically run out of memory, and require several minutes before it could compile again.

Any Questions?

swhite@naughtydog.com